

Denoising Diffusion Implicit Models

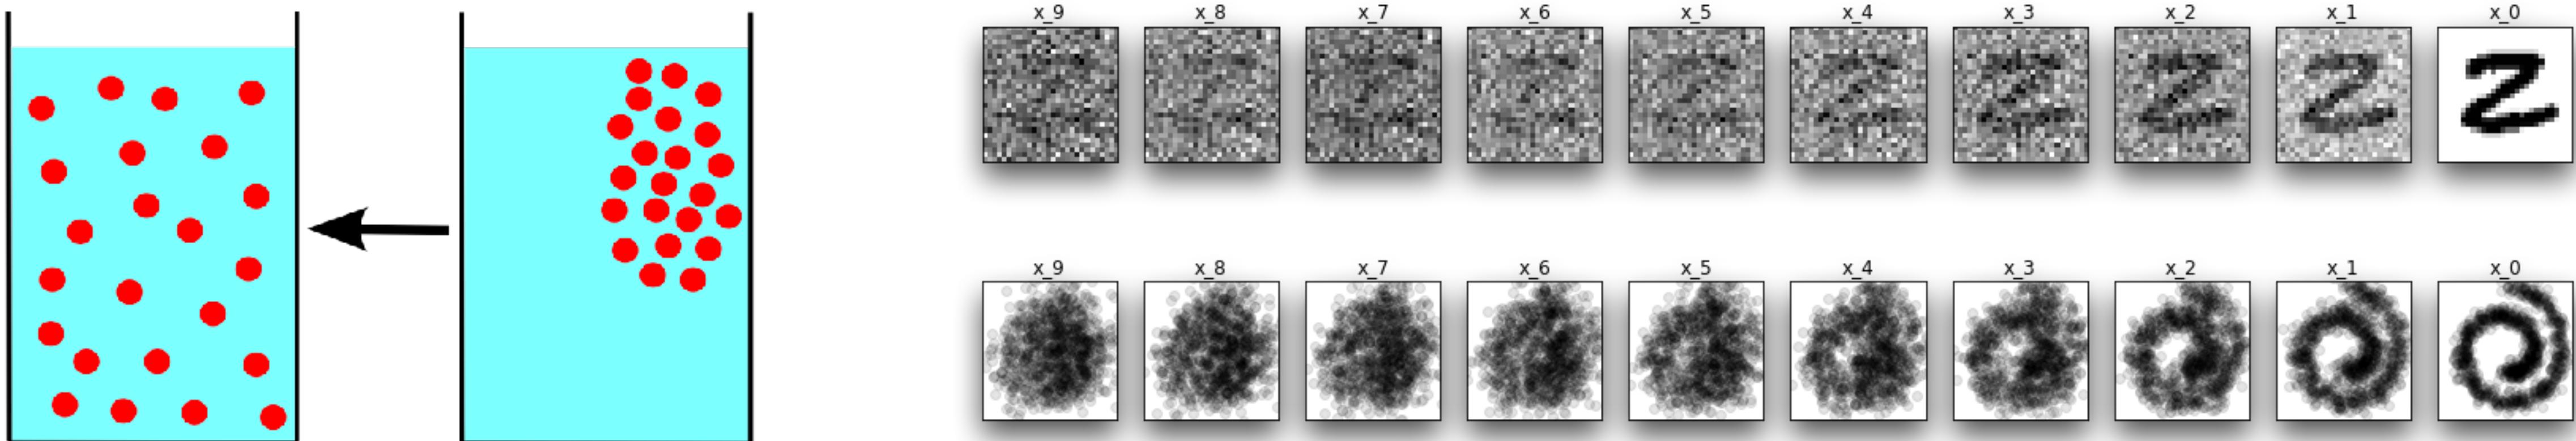
Learning with Keras Code Example

SCPARK

Diffusion Model

Diffusion Model

- Sohl-Dickstein, Jascha의 논문 Deep unsupervised learning using nonequilibrium thermodynamics에서 제안
- 물리에서의 diffusion 현상에 영감을 받아 모델을 설계

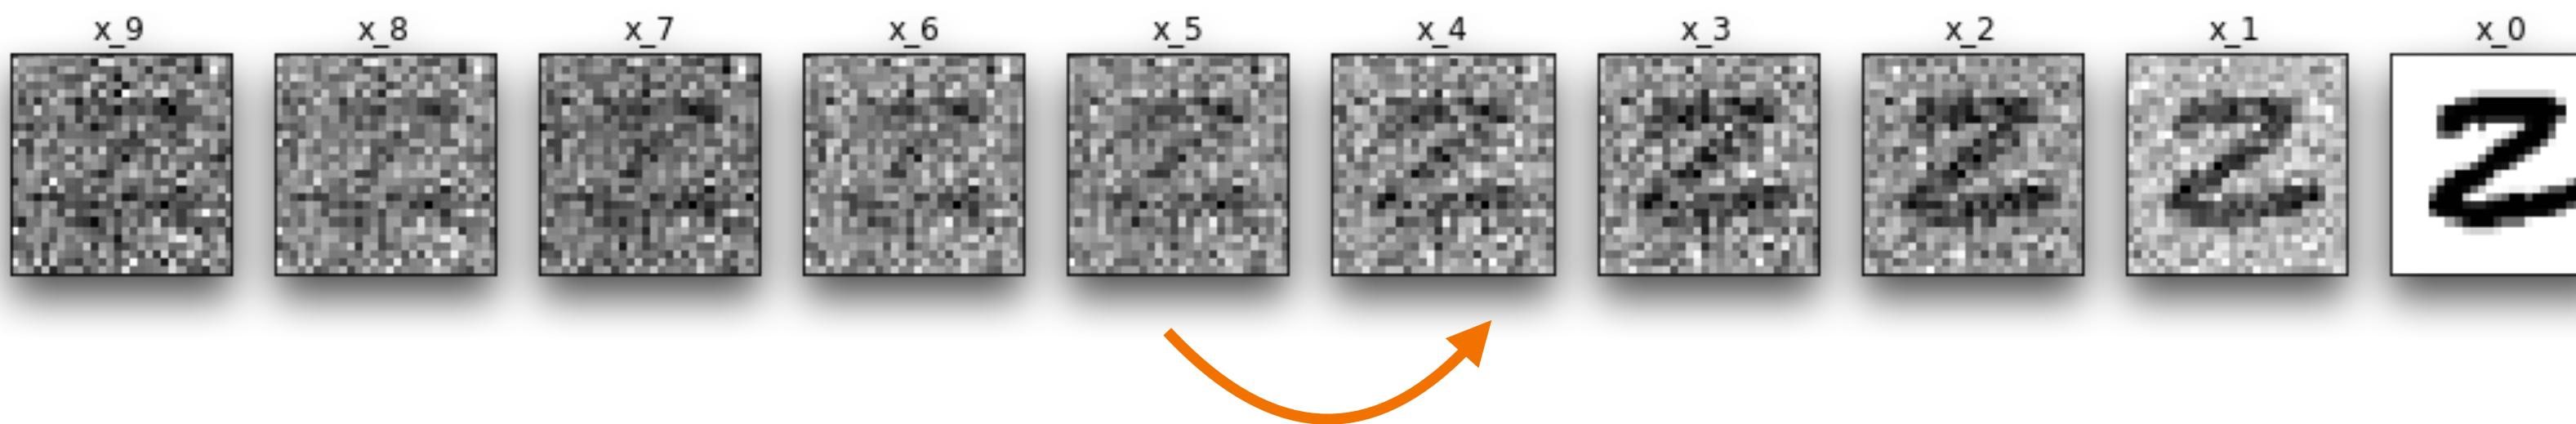


Diffusion

[https://en.wikipedia.org/wiki/Diffusion
\(Flipped\)](https://en.wikipedia.org/wiki/Diffusion_(Flipped))

Diffusion Model

Forward Process $q(\mathbf{x}_{1:T} \mid \mathbf{x}_0) := \prod_{t=1}^T q(\mathbf{x}_t \mid \mathbf{x}_{t-1}), q(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{x}_0) := \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$



Posterior $q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t \mathbf{I})$

where $\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \frac{\sqrt{\bar{\alpha}_{t-1}} \beta_t}{1 - \bar{\alpha}_t} \mathbf{x}_0 + \frac{\sqrt{\bar{\alpha}_t} (1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t$ and $\tilde{\beta}_t := \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t$

**Backward Process
(Neural Networks)**

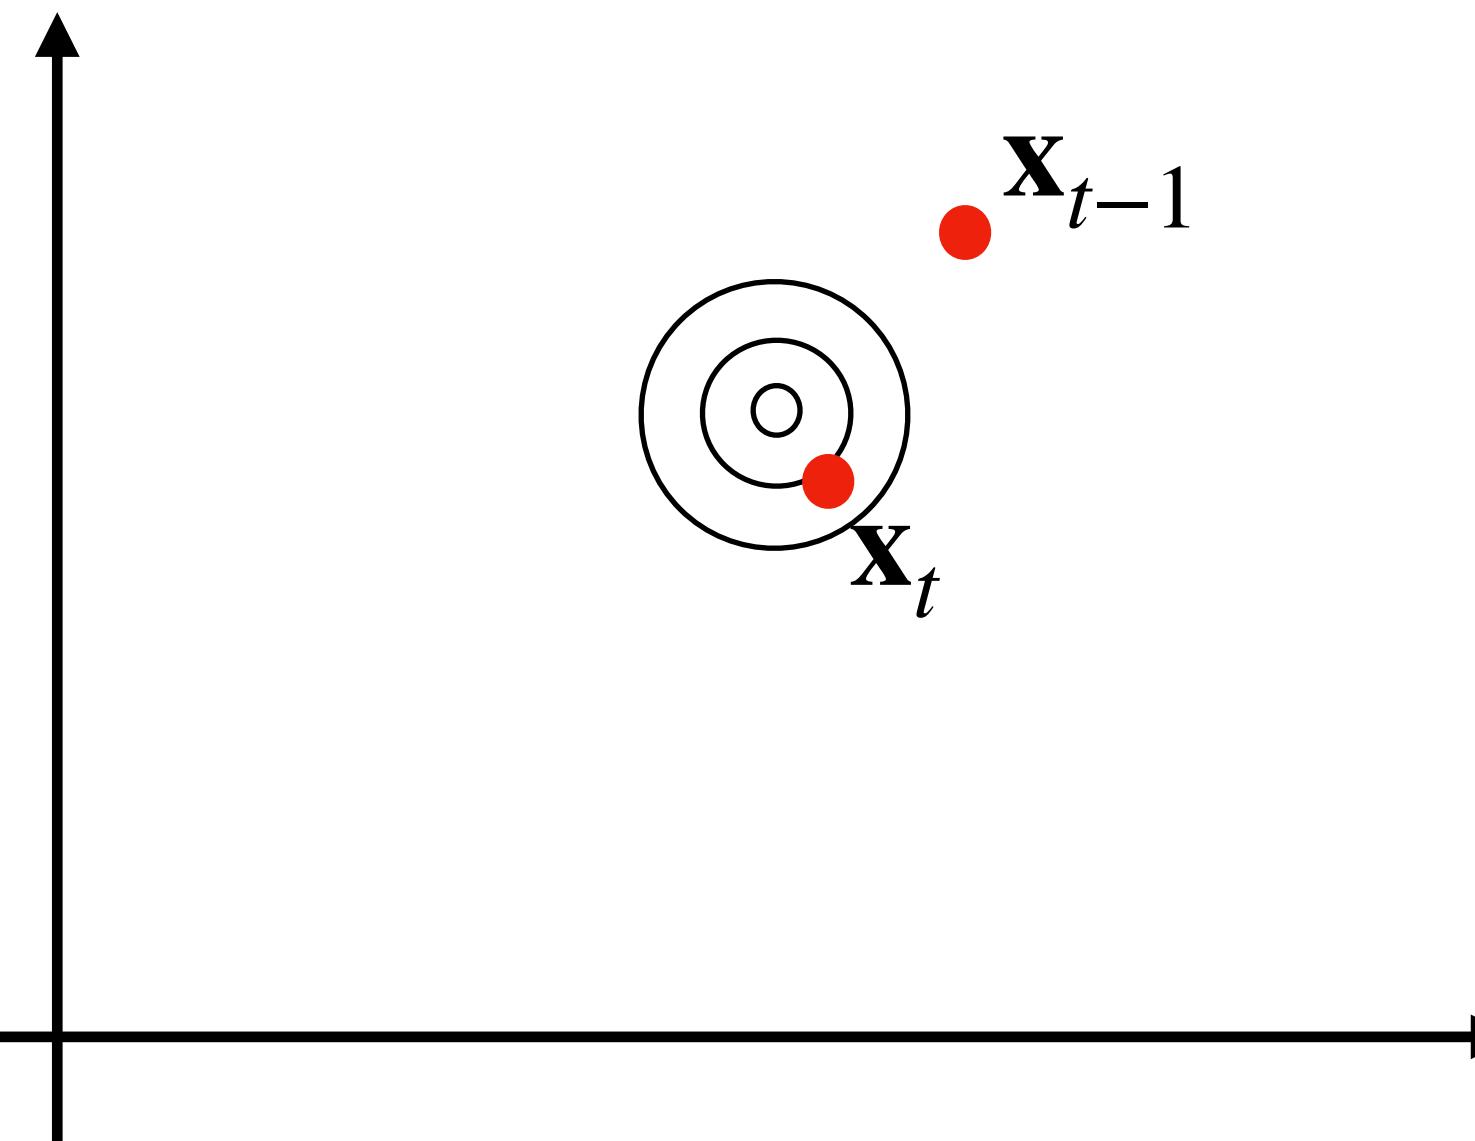
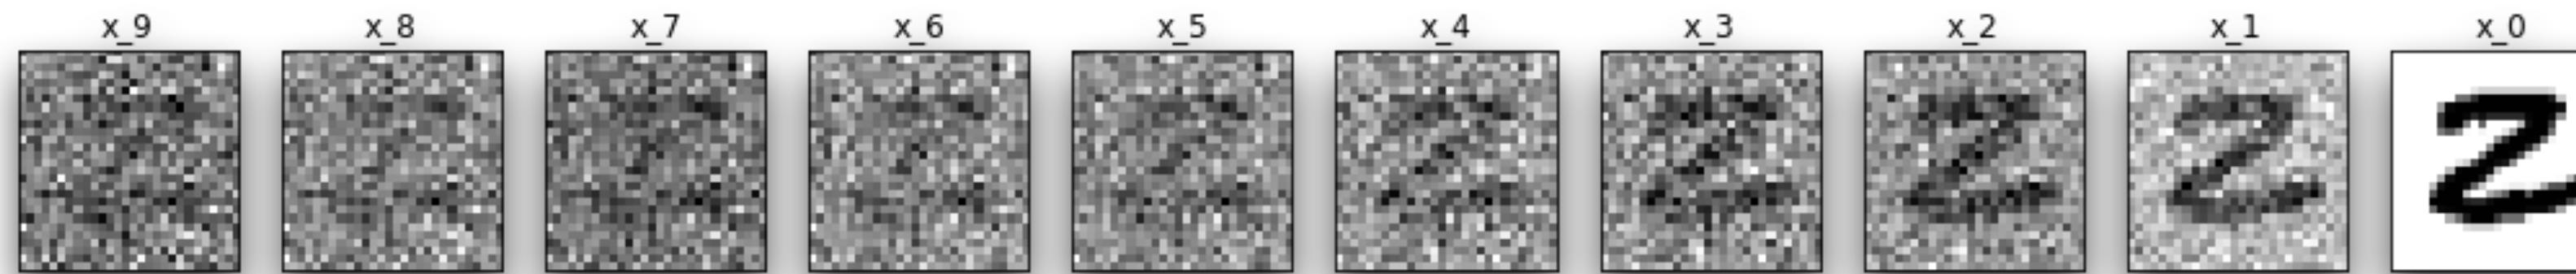
$p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t) := \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_\theta(\mathbf{x}_t, t), \boldsymbol{\Sigma}_\theta(\mathbf{x}_t, t))$

Loss Function

$$D_{\text{KL}}(q(\mathbf{x}_{t-1} \mid \mathbf{x}_t, \mathbf{x}_0) \parallel p_\theta(\mathbf{x}_{t-1} \mid \mathbf{x}_t))$$

Diffusion Model - Forward Process

Forward Process $q(\mathbf{x}_{1:T} \mid \mathbf{x}_0) := \prod_{t=1}^T q(\mathbf{x}_t \mid \mathbf{x}_{t-1}), q(\mathbf{x}_t \mid \mathbf{x}_{t-1}, \mathbf{x}_0) := \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$



Distribution of \mathbf{x}_t at an arbitrary timestep t in closed form

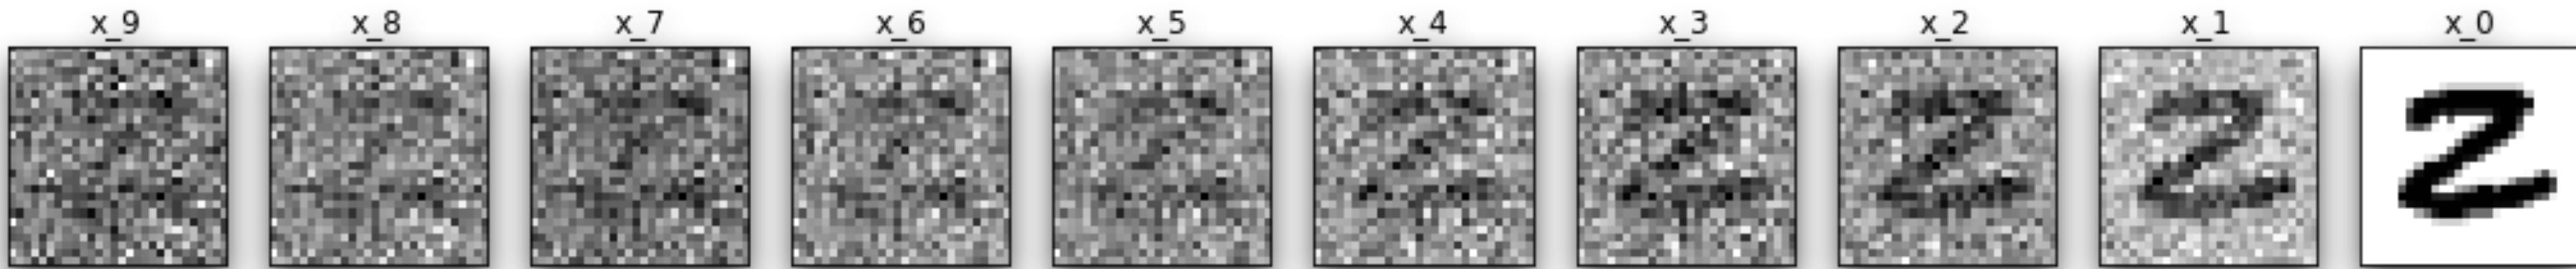
$$q(\mathbf{x}_t \mid \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I}), \text{ where } \alpha_t := 1 - \beta_t \text{ and } \bar{\alpha}_t := \prod_{s=1}^t \alpha_s$$

식 유도는 Lil'Log 참고

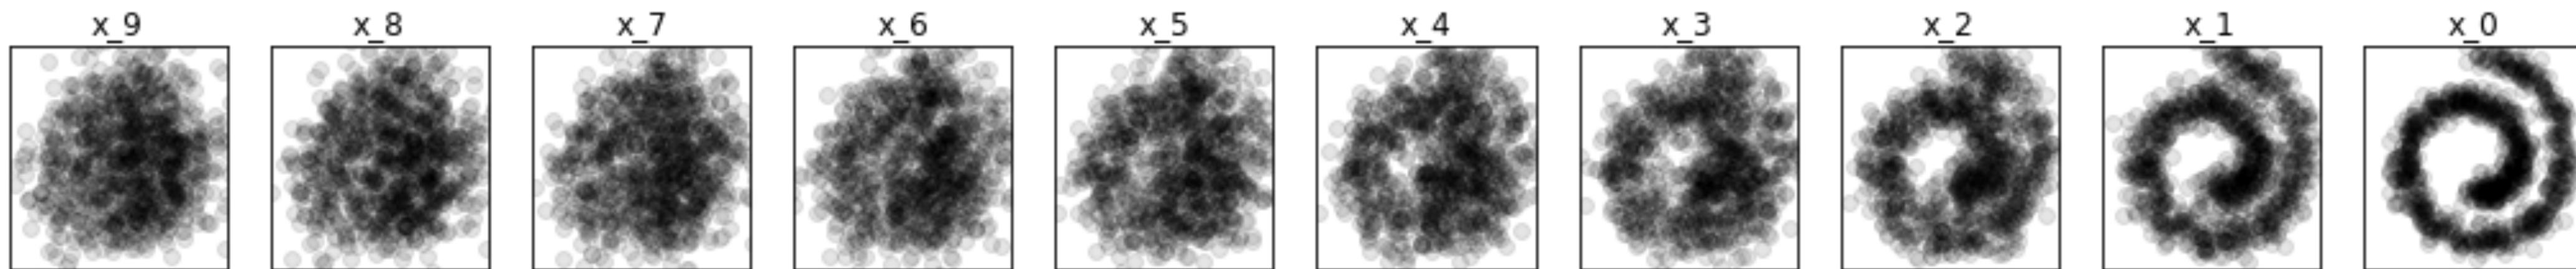
<https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>

Diffusion Model - Forward Process

MNIST single data, $\beta = 0.2$, $T = 10$

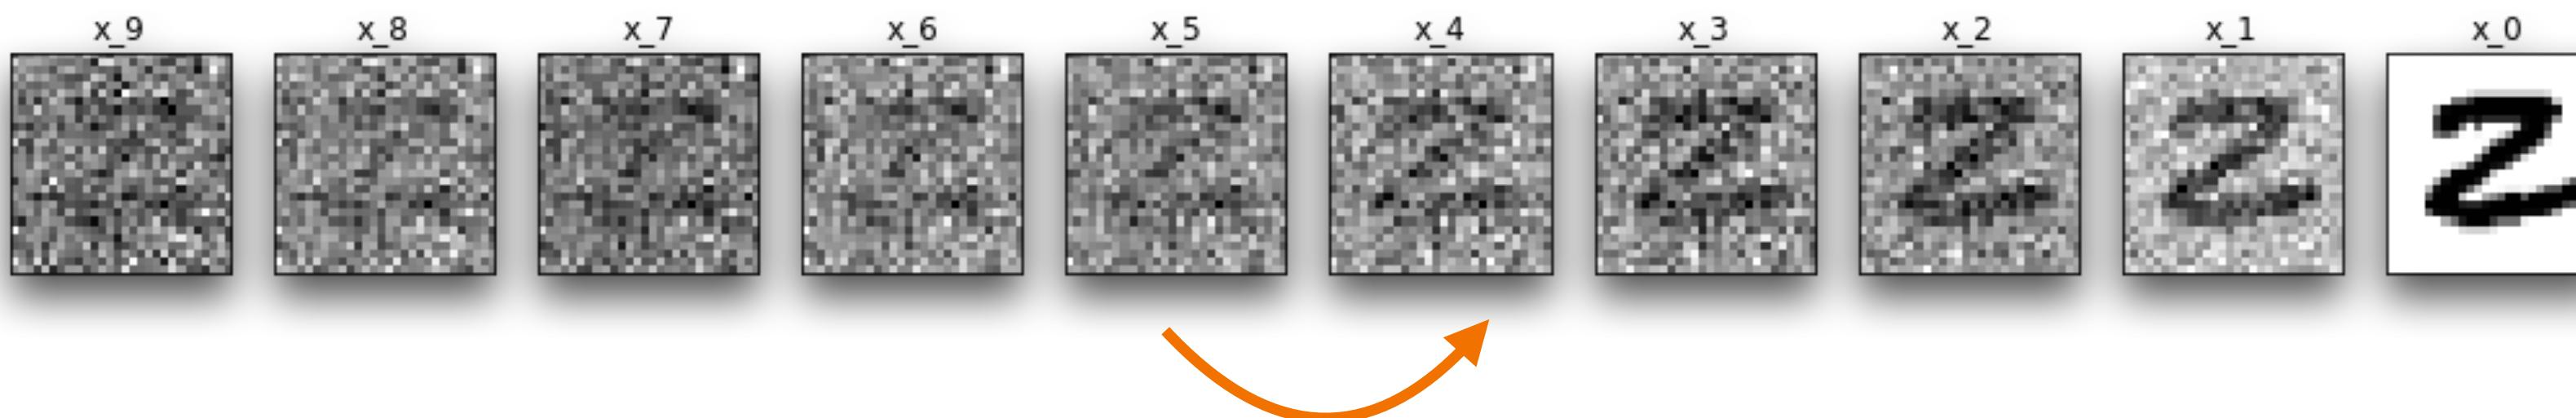


Swiss roll dataset, $\beta = 0.05$, $T = 10$



Diffusion Model - Posterior

Forward Process $q(\mathbf{x}_{1:T} | \mathbf{x}_0) := \prod_{t=1}^T q(\mathbf{x}_t | \mathbf{x}_{t-1})$, $q(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{x}_0) := \mathcal{N}(\mathbf{x}_t; \sqrt{1-\beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I})$



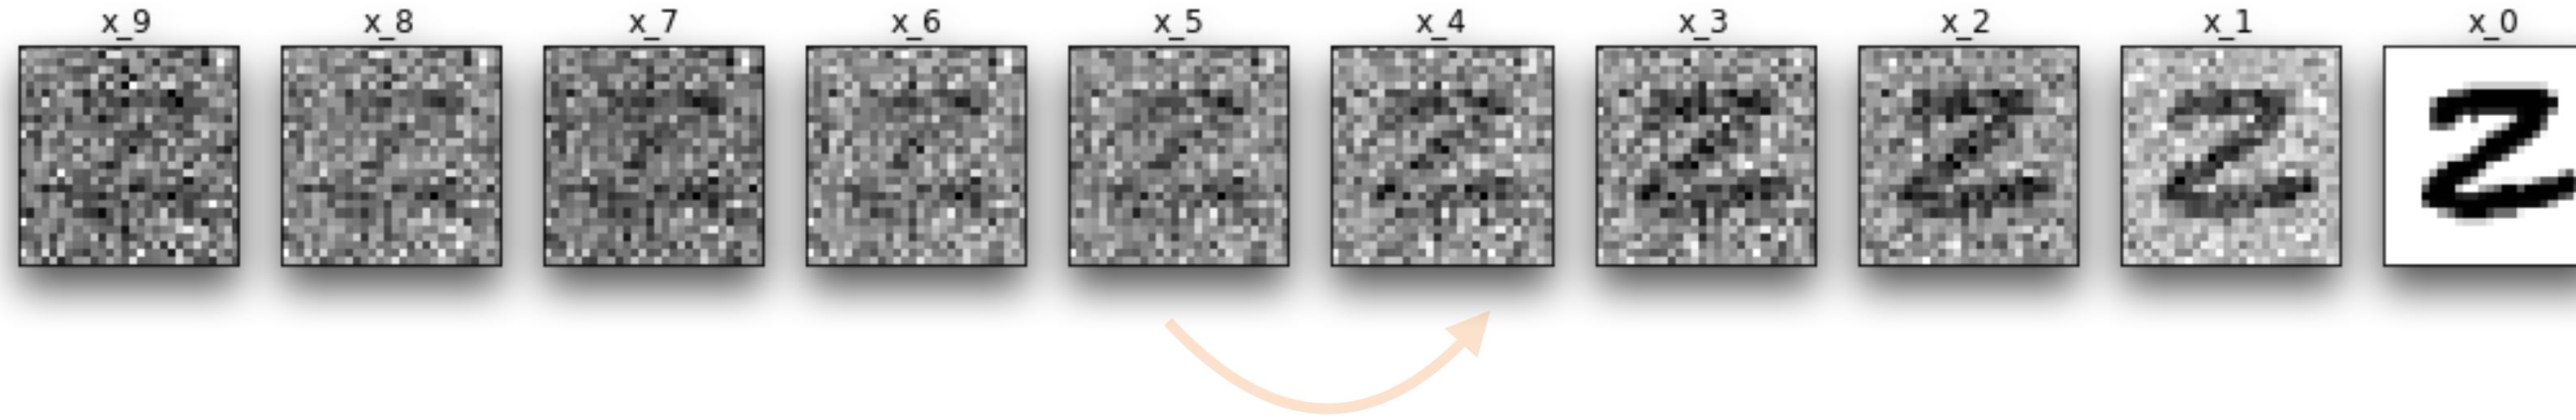
Posterior $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t\mathbf{I})$

$$\text{where } \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1-\bar{\alpha}_t}\mathbf{x}_0 + \frac{\sqrt{\alpha_t}(1-\bar{\alpha}_{t-1})}{1-\bar{\alpha}_t}\mathbf{x}_t \quad \text{and} \quad \tilde{\beta}_t := \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t}\beta_t$$

$$q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \frac{q(\mathbf{x}_{t-1} | \mathbf{x}_0)q(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{x}_0)}{q(\mathbf{x}_t | \mathbf{x}_0)} \text{ by Bayes' Rule}$$

Diffusion Model - Backward Process

Forward Process $q(\mathbf{x}_{1:T} | \mathbf{x}_0) := \prod_{t=1}^T q(\mathbf{x}_t | \mathbf{x}_{t-1}), q(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{x}_0) := \mathcal{N}(\mathbf{x}_t; \sqrt{1-\beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I})$



Posterior $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t\mathbf{I})$

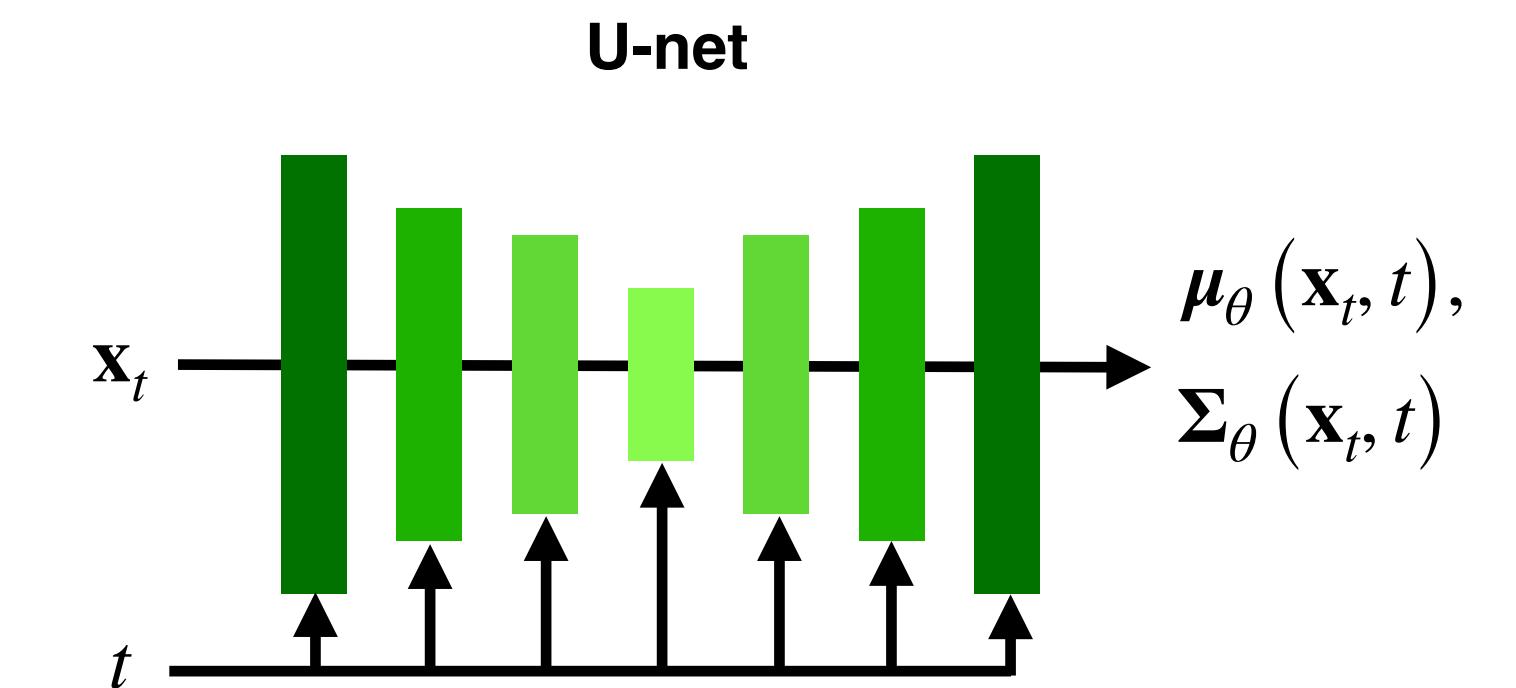
$$\text{where } \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1-\bar{\alpha}_t}\mathbf{x}_0 + \frac{\sqrt{\bar{\alpha}_t}(1-\bar{\alpha}_{t-1})}{1-\bar{\alpha}_t}\mathbf{x}_t \quad \text{and} \quad \tilde{\beta}_t := \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t}\beta_t$$

**Backward Process
(Neural Networks)**

$p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) := \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_\theta(\mathbf{x}_t, t), \boldsymbol{\Sigma}_\theta(\mathbf{x}_t, t))$

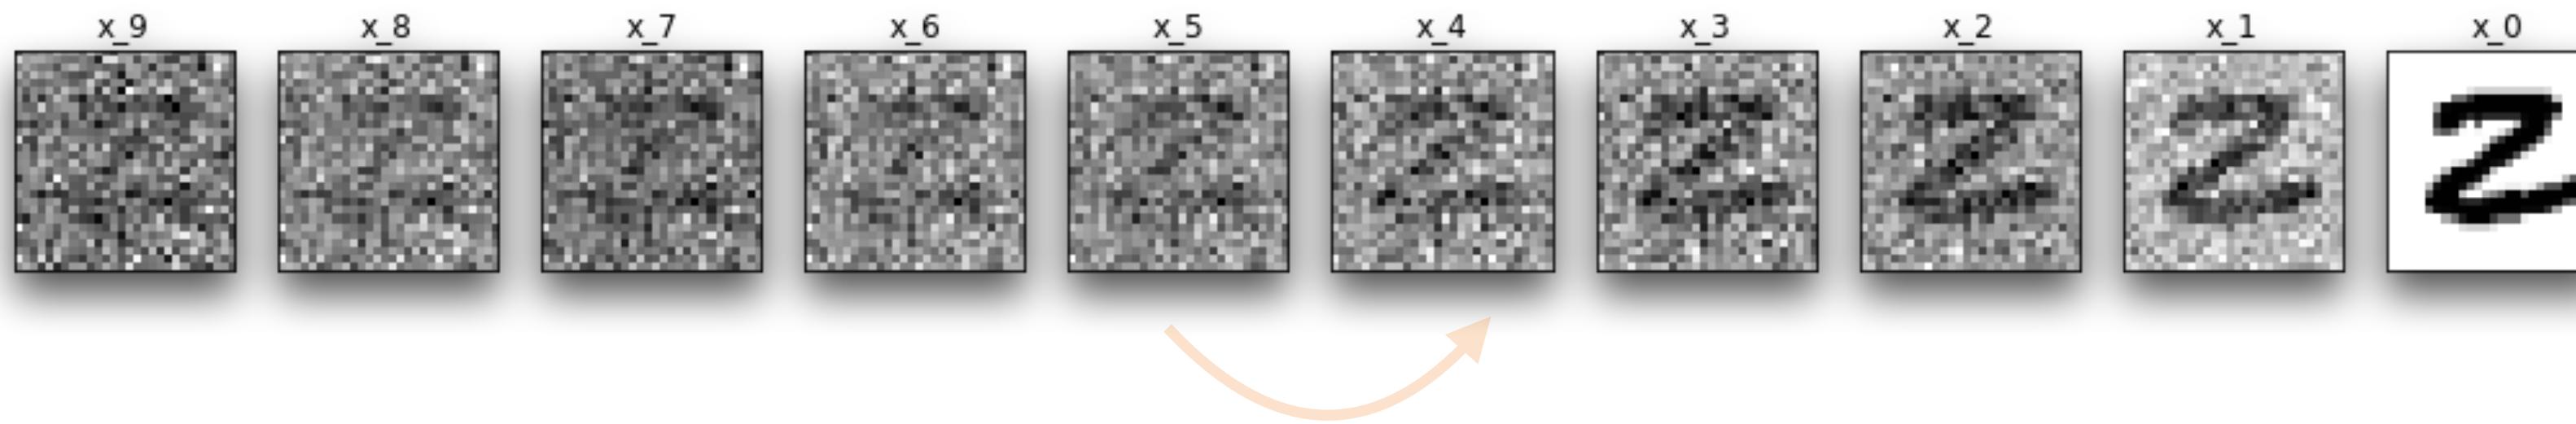
Loss Function

$$D_{\text{KL}}(q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) \| p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t))$$



Diffusion Model - Loss Function

Forward Process $q(\mathbf{x}_{1:T} | \mathbf{x}_0) := \prod_{t=1}^T q(\mathbf{x}_t | \mathbf{x}_{t-1})$, $q(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{x}_0) := \mathcal{N}(\mathbf{x}_t; \sqrt{1-\beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I})$



Posterior $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t\mathbf{I})$

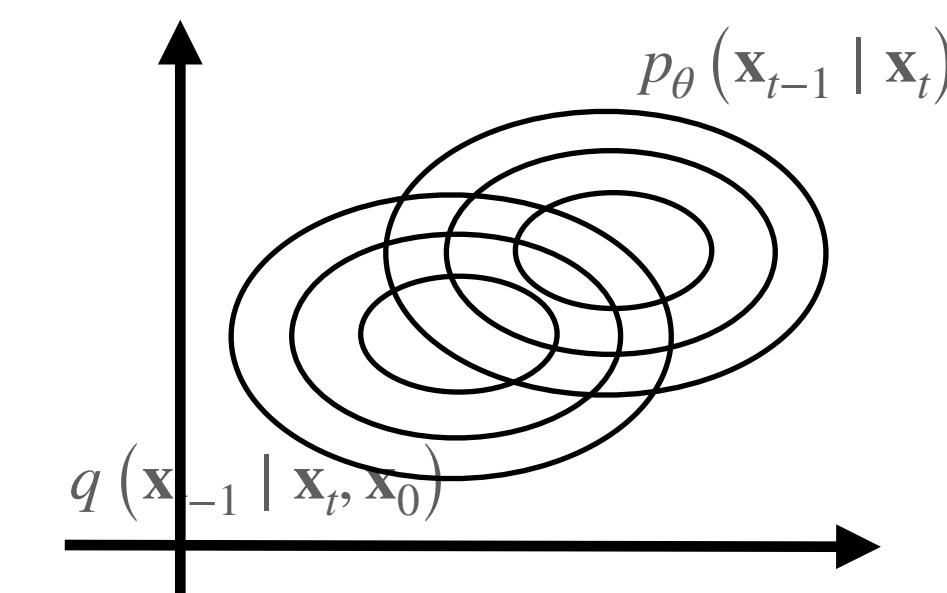
$$\text{where } \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1-\bar{\alpha}_t}\mathbf{x}_0 + \frac{\sqrt{\bar{\alpha}_t}(1-\bar{\alpha}_{t-1})}{1-\bar{\alpha}_t}\mathbf{x}_t \quad \text{and} \quad \tilde{\beta}_t := \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t}\beta_t$$

**Backward Process
(Neural Networks)**

$p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) := \mathcal{N}(\mathbf{x}_{t-1}; \mu_\theta(\mathbf{x}_t, t), \Sigma_\theta(\mathbf{x}_t, t))$

Loss Function

$$D_{\text{KL}}(q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) \| p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t))$$



Diffusion Model - Output Samples

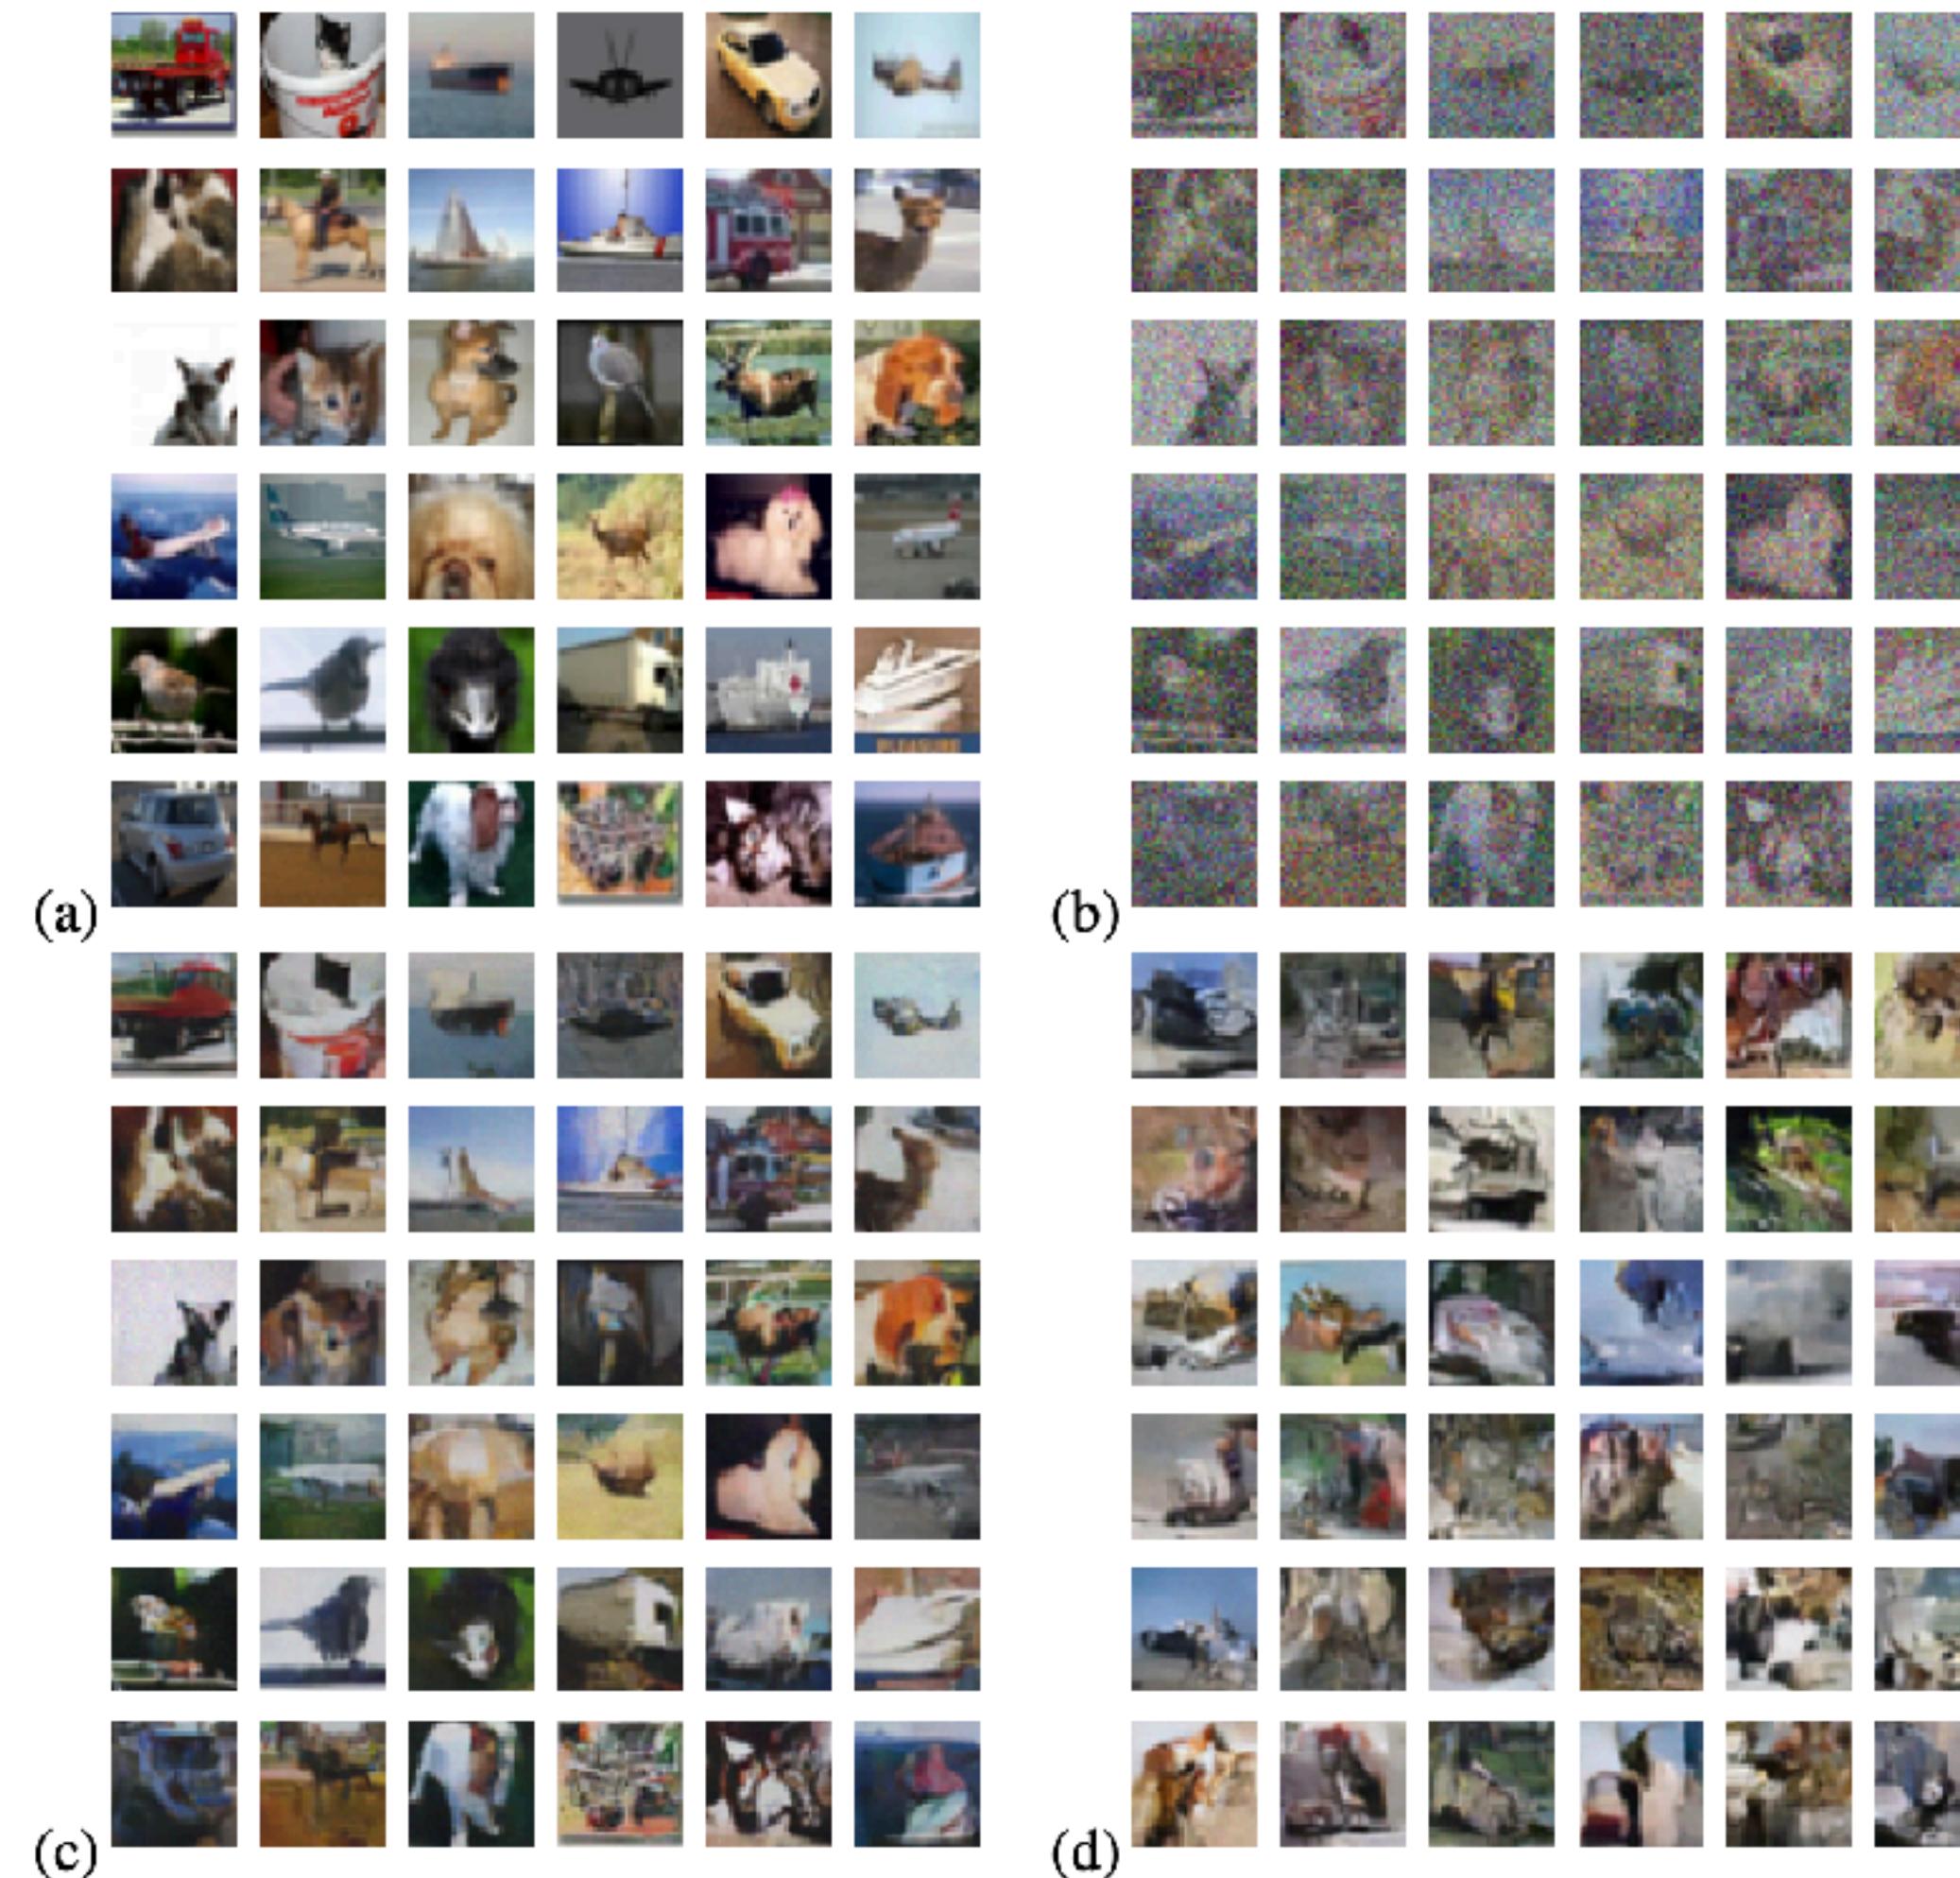


Figure 3. The proposed framework trained on the CIFAR-10 ([Krizhevsky & Hinton, 2009](#)) dataset. (a) Example holdout data (similar to training data). (b) Holdout data corrupted with Gaussian noise of variance 1 (SNR = 1). (c) Denoised images, generated by sampling from the posterior distribution over denoised images conditioned on the images in (b). (d) Samples generated by the diffusion model.

DDPM

DDPM (Denoising Diffusion Probabilistic Models)

- Jonathan Ho의 논문 Denoising diffusion probabilistic models에서 제안

Posterior $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}\left(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t \mathbf{I}\right)$

Distribution of \mathbf{x}_t at an arbitrary timestep t in closed form

$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}\left(\mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I}\right), \text{ where } \alpha_t := 1 - \beta_t \text{ and } \bar{\alpha}_t := \prod_{s=1}^t \alpha_s$$

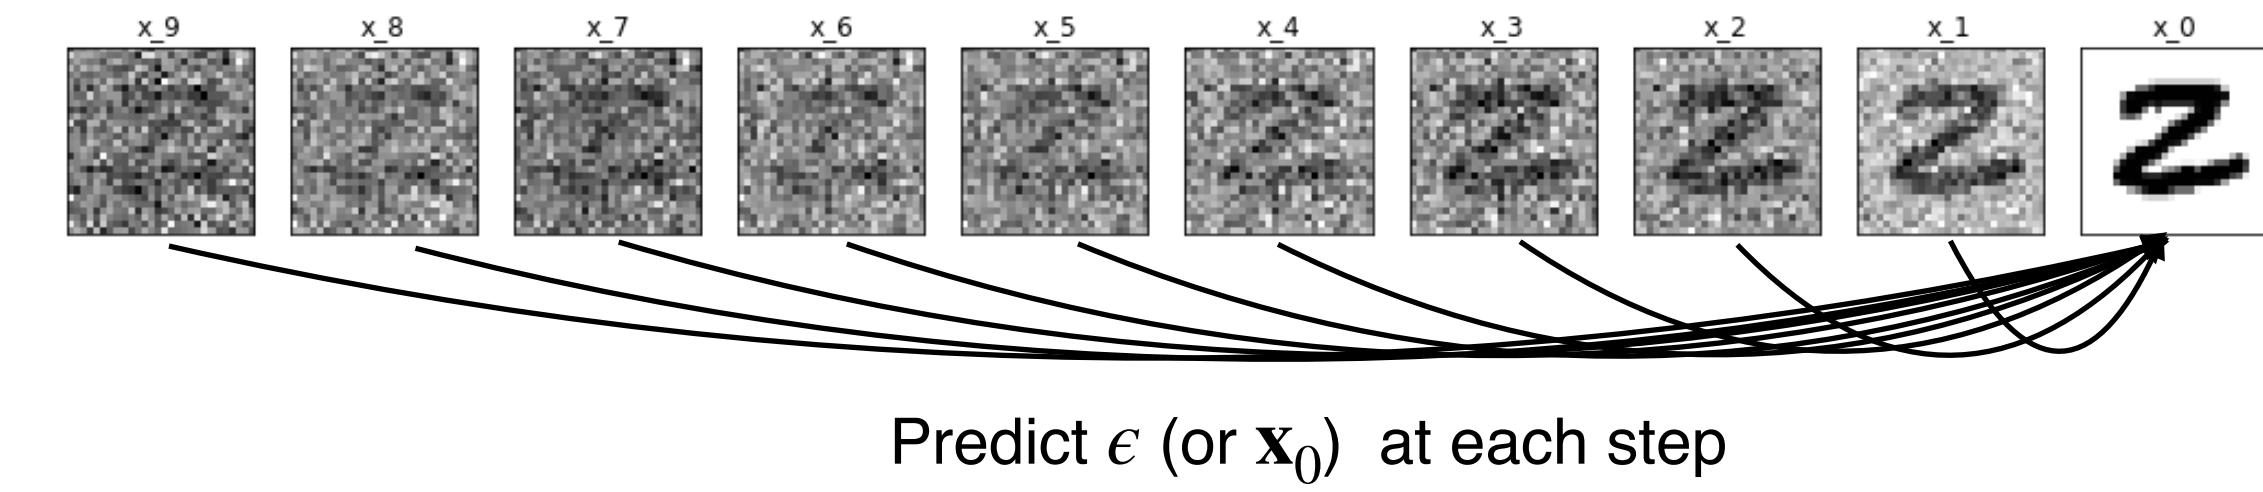
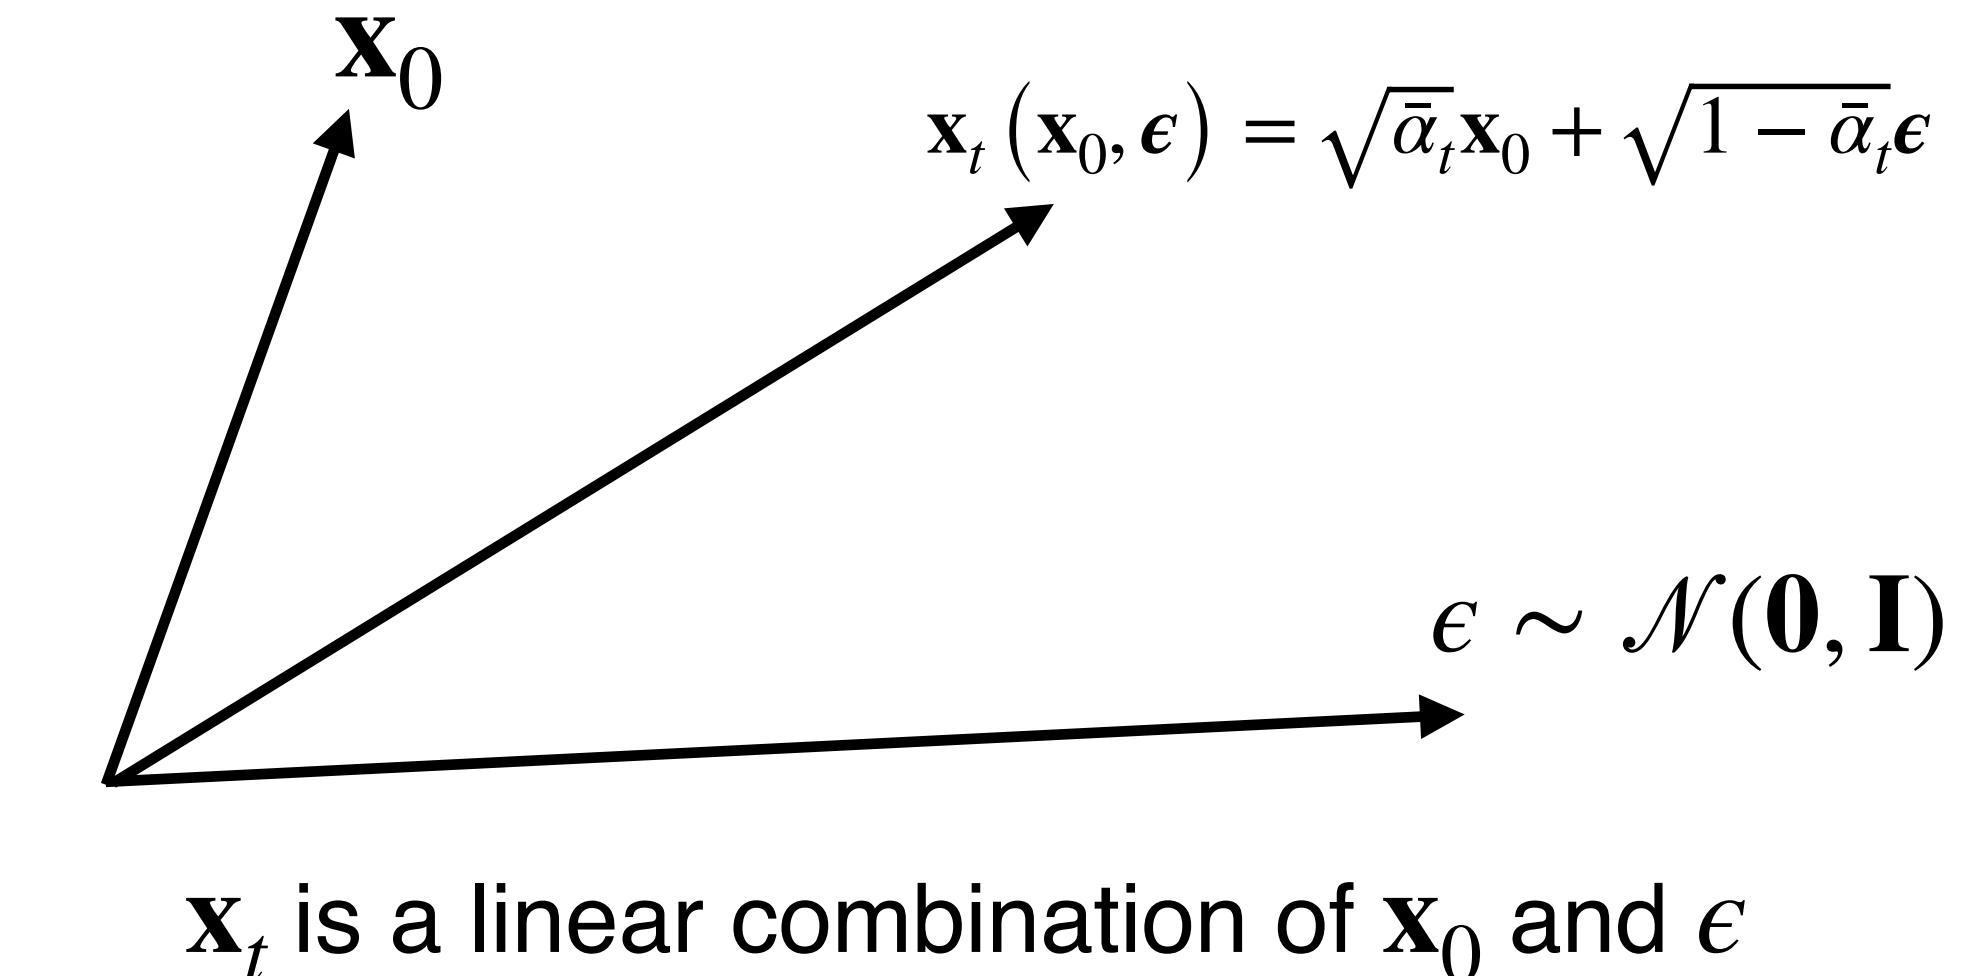
Loss Function $L_{\text{simple}}(\theta) := \mathbb{E}_{t, \mathbf{x}_0, \epsilon} \left[\left\| \epsilon - \epsilon_\theta \left(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t \right) \right\|^2 \right]$

```

106     def p_loss(self, model, x_0, t, noise=None):
107         if noise is None:
108             noise = torch.randn_like(x_0)           generate ε
109
110             x_noise = self.q_sample(x_0, t, noise) sample x_t
111             x_recon = model(x_noise, t)           predict ε
112
113             return F.mse_loss(x_recon, noise)

```

<https://github.com/robinly/denoising-diffusion-pytorch/blob/master/diffusion.py>



DDPM - Output Samples

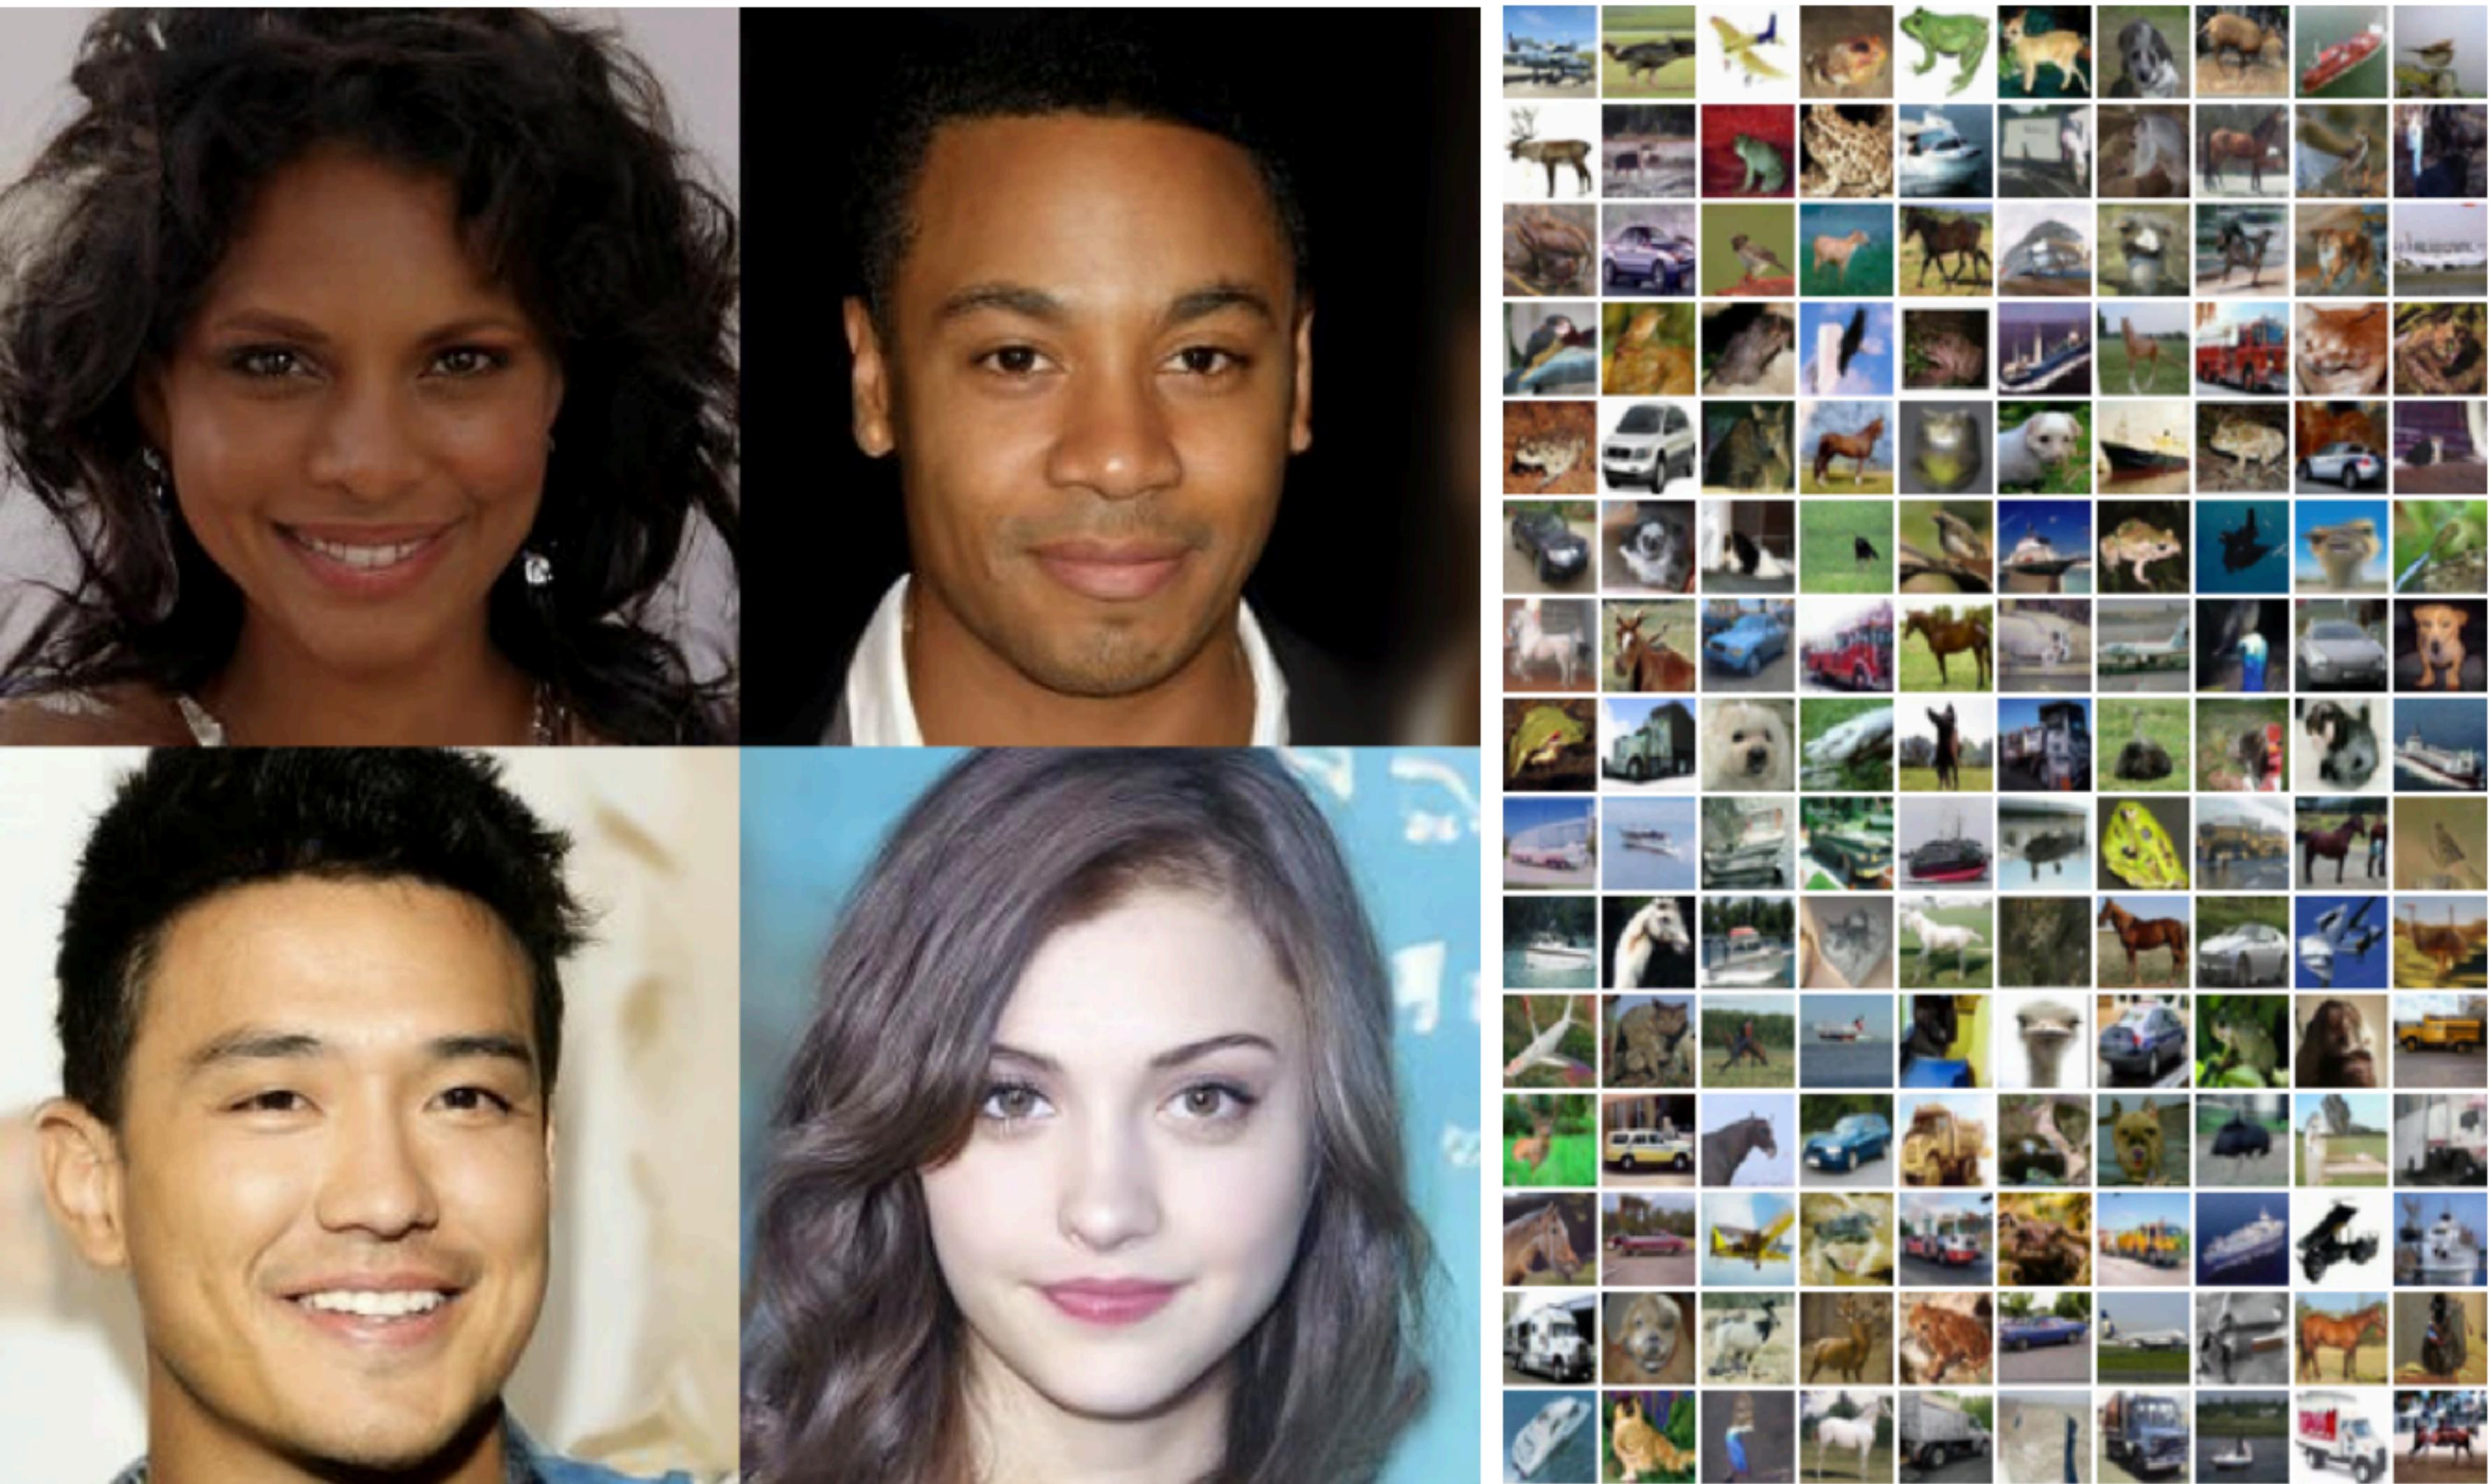
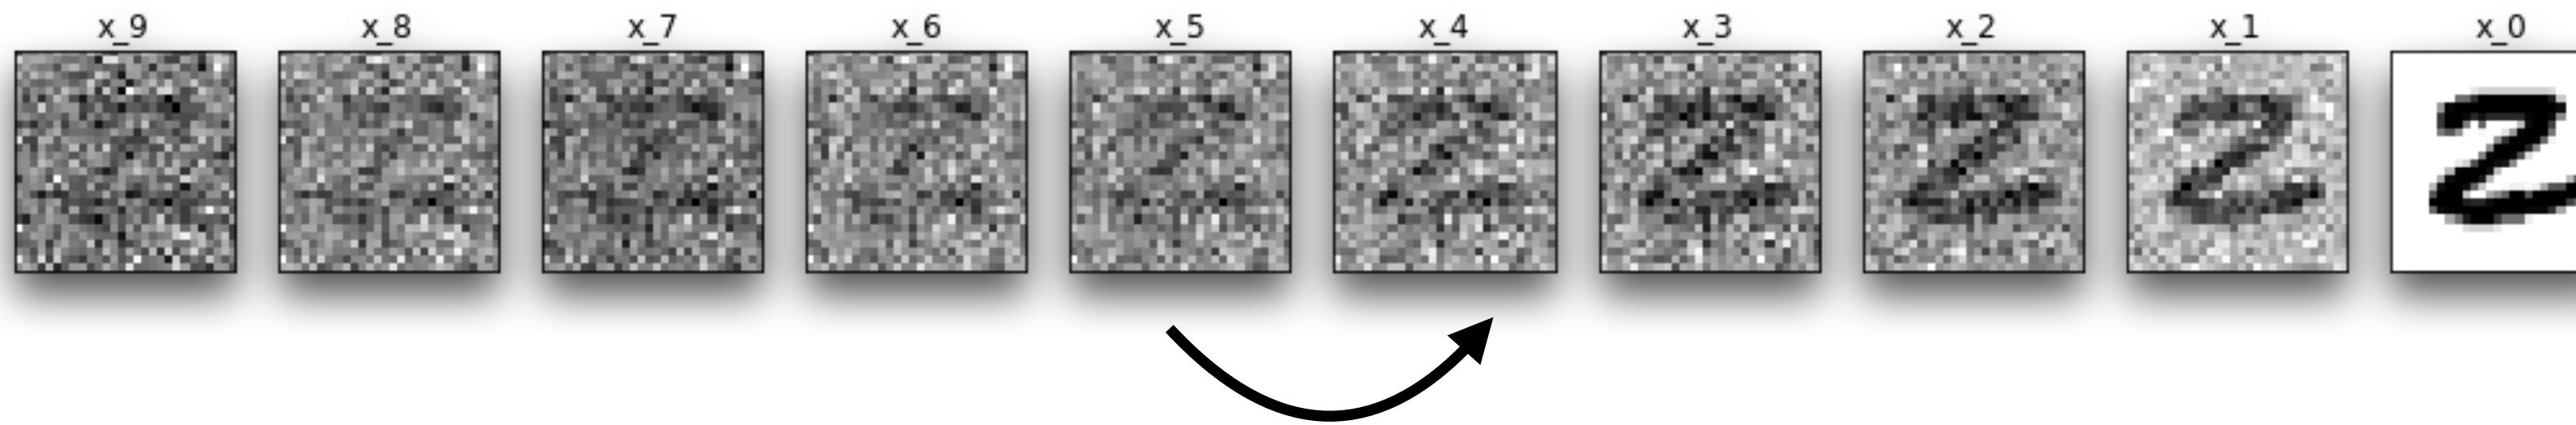


Figure 1: Generated samples on CelebA-HQ 256×256 (left) and unconditional CIFAR10 (right)

DDPM의 문제점 : 너무 느리다!!!

- DDPM의 posterior는 noised image \mathbf{x}_t 와 denoised image \mathbf{x}_0 가 주어졌을 때, 이전 step의 noised image \mathbf{x}_{t-1} 를 예측



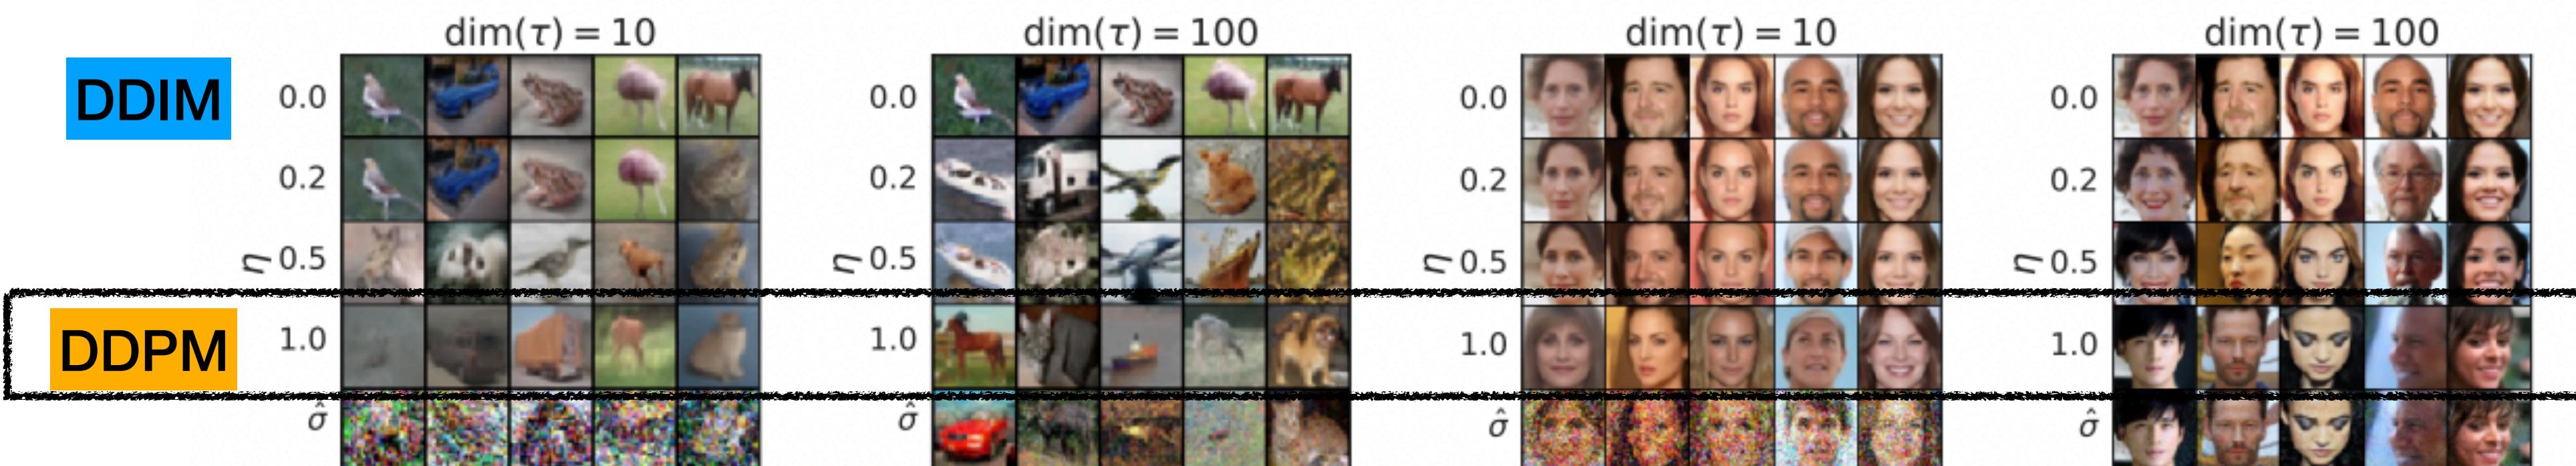
DDPM Posterior $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}\left(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t \mathbf{I}\right)$

where $\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1-\bar{\alpha}_t}\mathbf{x}_0 + \frac{\sqrt{\alpha_t}(1-\bar{\alpha}_{t-1})}{1-\bar{\alpha}_t}\mathbf{x}_t$ and $\tilde{\beta}_t := \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t}\beta_t$

DDPM의 문제점 : 너무 느리다!!!

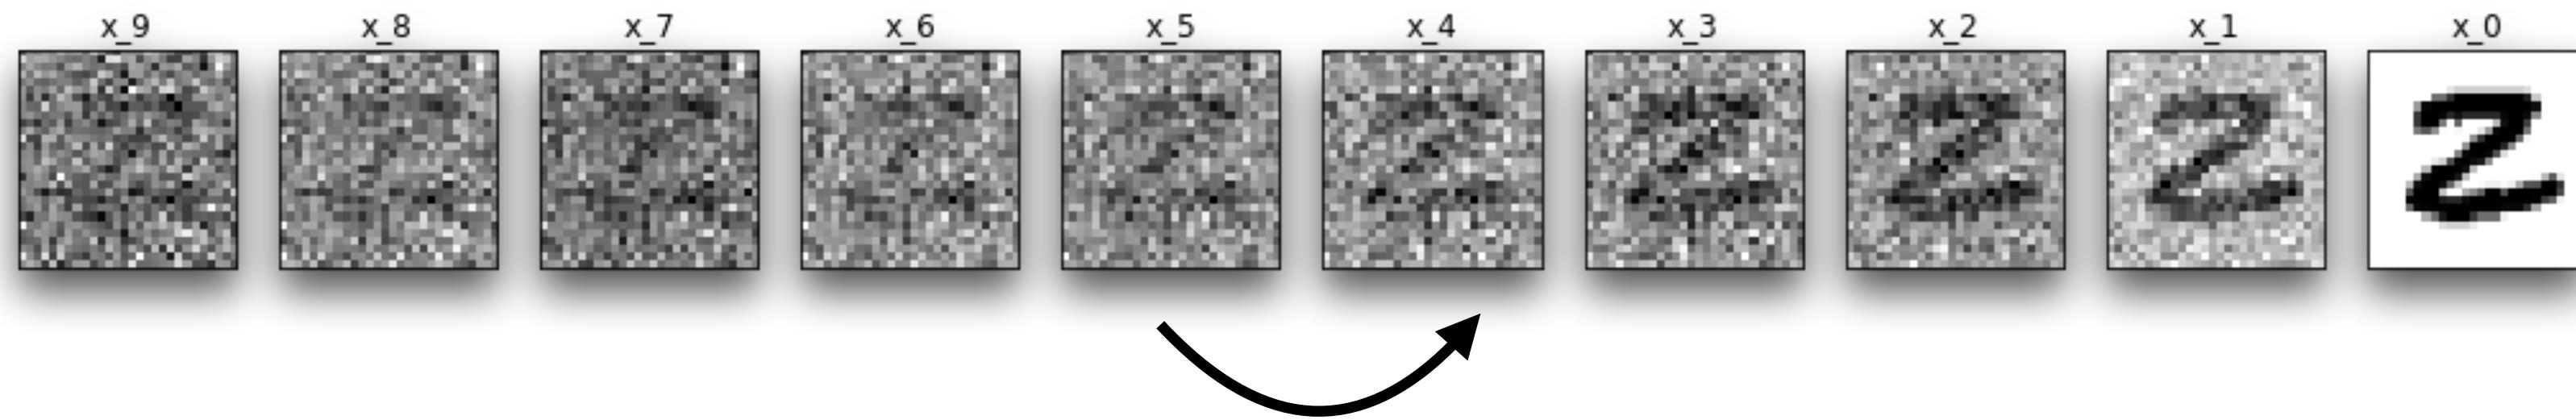
- 10steps, 20step, 50steps, 100steps 씩 띄엄띄엄 예측해나가면 어떨까?
- 가능하긴 하지만 sampling하는 총 step 수가 적어짐에 따라 심각하게 quality가 저하됨

S	CIFAR10 (32×32)					CelebA (64×64)				
	10	20	50	100	1000	10	20	50	100	1000
DDIM	0.0	13.36	6.84	4.67	4.16	4.04	17.33	13.73	9.17	6.53
	0.2	14.04	7.11	4.77	4.25	4.09	17.66	14.11	9.51	6.79
	0.5	16.66	8.35	5.25	4.46	4.29	19.86	16.06	11.01	8.09
DDPM	1.0	41.07	18.36	8.01	5.78	4.73	33.12	26.03	18.48	13.93
$\hat{\sigma}$		367.43	133.37	32.72	9.99	3.17	299.71	183.83	71.71	45.20
										3.26



DDIM

- Noised image \mathbf{x}_t 와 denoised image \mathbf{x}_0 가 주어졌을 때, \mathbf{x}_{t-1} 를 예측하는 posterior부터 다시 설계해보자!
(DDPM은 forward process부터 설계하고 posterior를 유도해 냄)



DDPM Posterior $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}\left(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t \mathbf{I}\right)$

where $\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1-\bar{\alpha}_t}\mathbf{x}_0 + \frac{\sqrt{\alpha_t}(1-\bar{\alpha}_{t-1})}{1-\bar{\alpha}_t}\mathbf{x}_t$ and $\tilde{\beta}_t := \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t}\beta_t$

Generalized Posterior $q_\sigma(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}\left(\sqrt{\alpha_{t-1}}\mathbf{x}_0 + \sqrt{1-\alpha_{t-1}-\sigma_t^2} \cdot \frac{\mathbf{x}_t - \sqrt{\alpha_t}\mathbf{x}_0}{\sqrt{1-\alpha_t}}, \sigma_t^2 \mathbf{I}\right)$

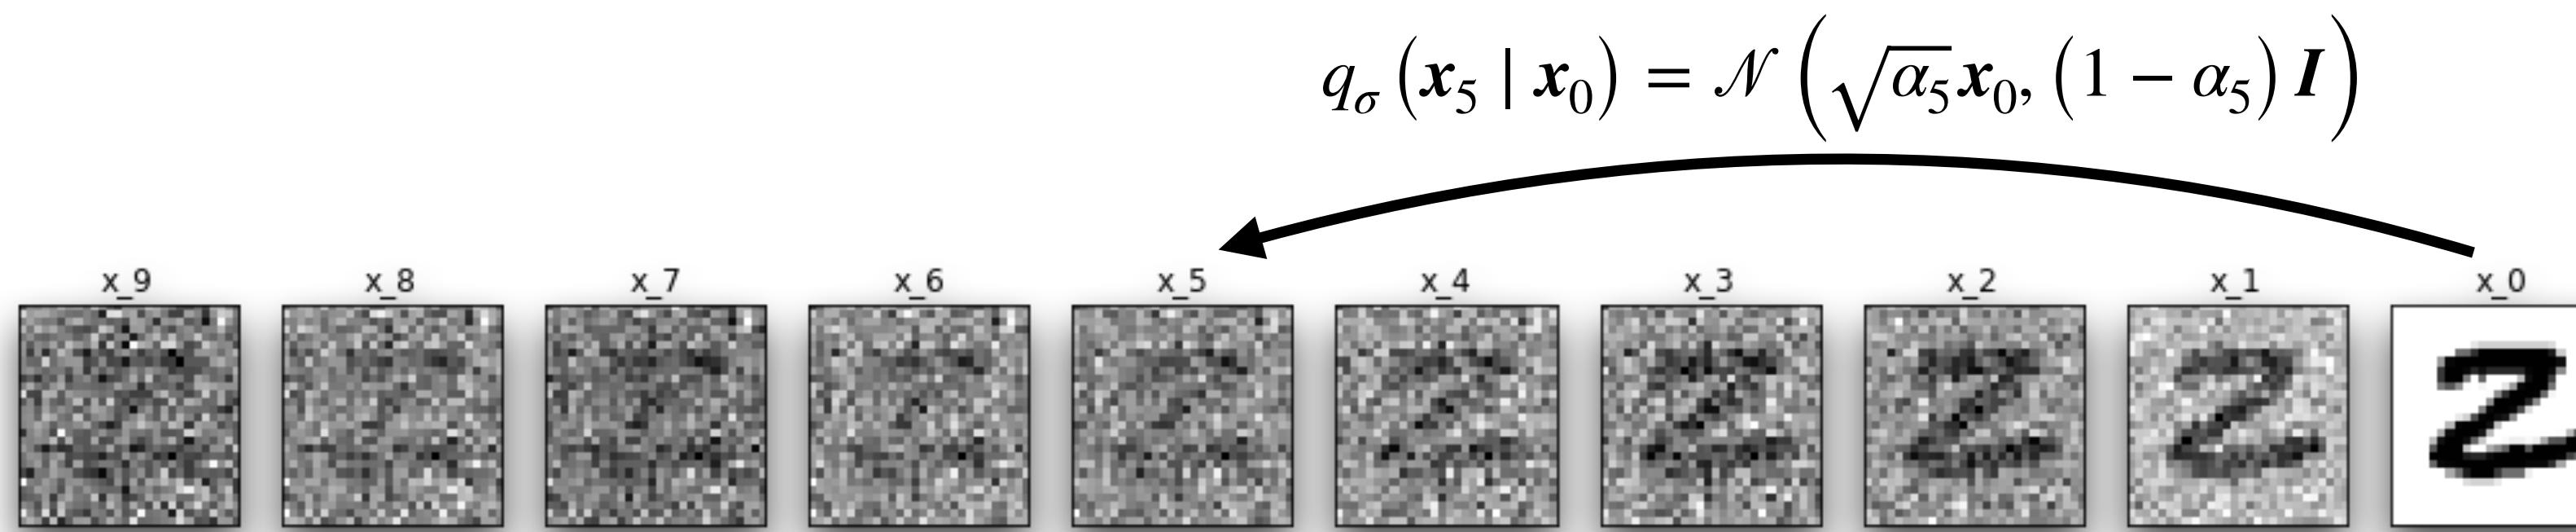
* DDPM에서 보이는 $\bar{\alpha}$ 는 DDIM에서 α 로 쓰임

DDIM

- Generalized posterior $q_\sigma(x_{t-1} | x_t, x_0)$ 는 어떻게 설계된 식인가?

>> DDPM에서 언급했던 $q_\sigma(x_t | x_0) = \mathcal{N}\left(\sqrt{\alpha_t}x_0, (1 - \alpha_t)\mathbf{I}\right)$ 식을 만족하도록 설계

>> 자세한 식 유도는 DDIM 논문 Appendix B에 언급

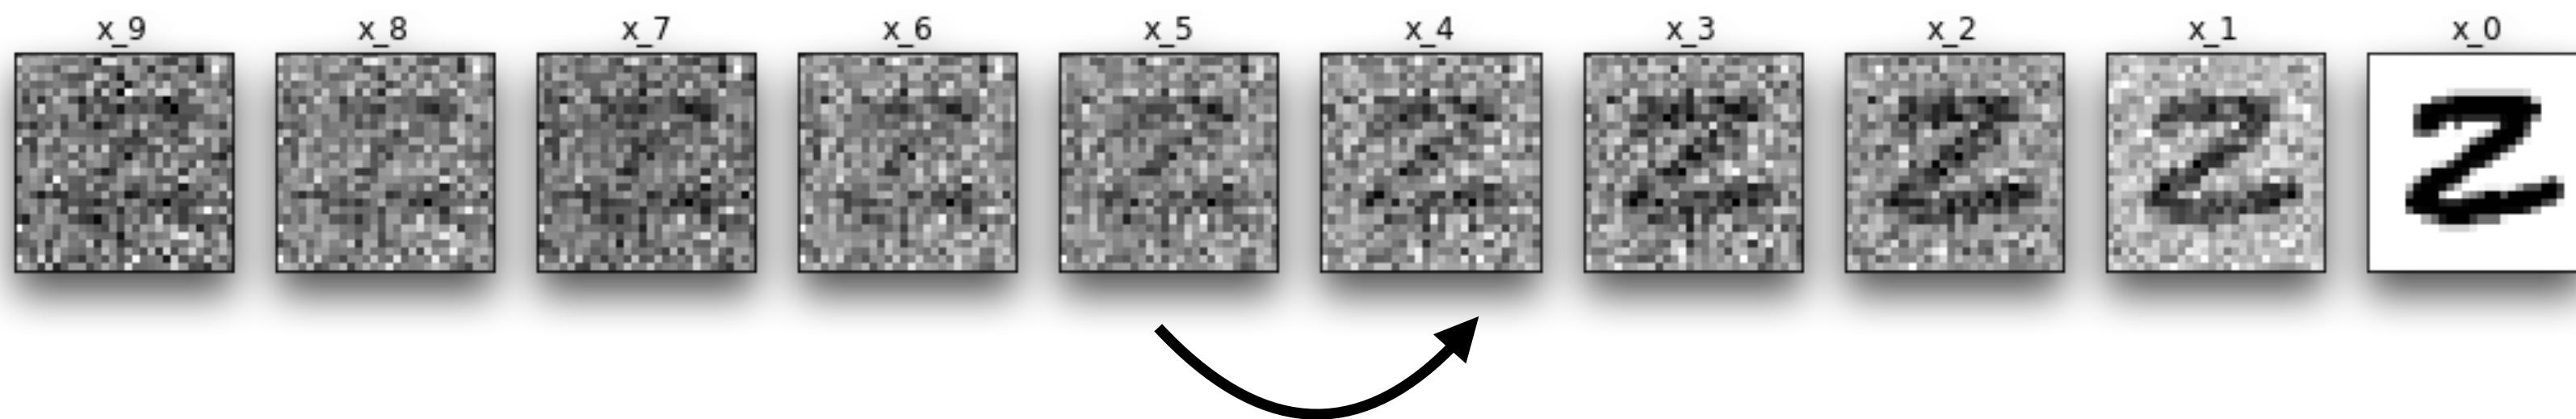


Distribution of x_t at an arbitrary timestep t in closed form

$$q_\sigma(x_t | x_0) = \mathcal{N}\left(\sqrt{\alpha_t}x_0, (1 - \alpha_t)\mathbf{I}\right)$$

DDIM

- Generalized posterior에서 $\sigma_t = \sqrt{(1 - \alpha_{t-1}) / (1 - \alpha_t)} \sqrt{1 - \alpha_t / \alpha_{t-1}}$ 으로 두면 DDPM과 같아짐
- $\sigma_t = 0$ 으로 두면 posterior가 deterministic function이 되고 이것을 DDIM이라고 부름



DDPM Posterior $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}\left(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t \mathbf{I}\right)$

where $\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t} \mathbf{x}_0 + \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t$ and $\tilde{\beta}_t := \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t$

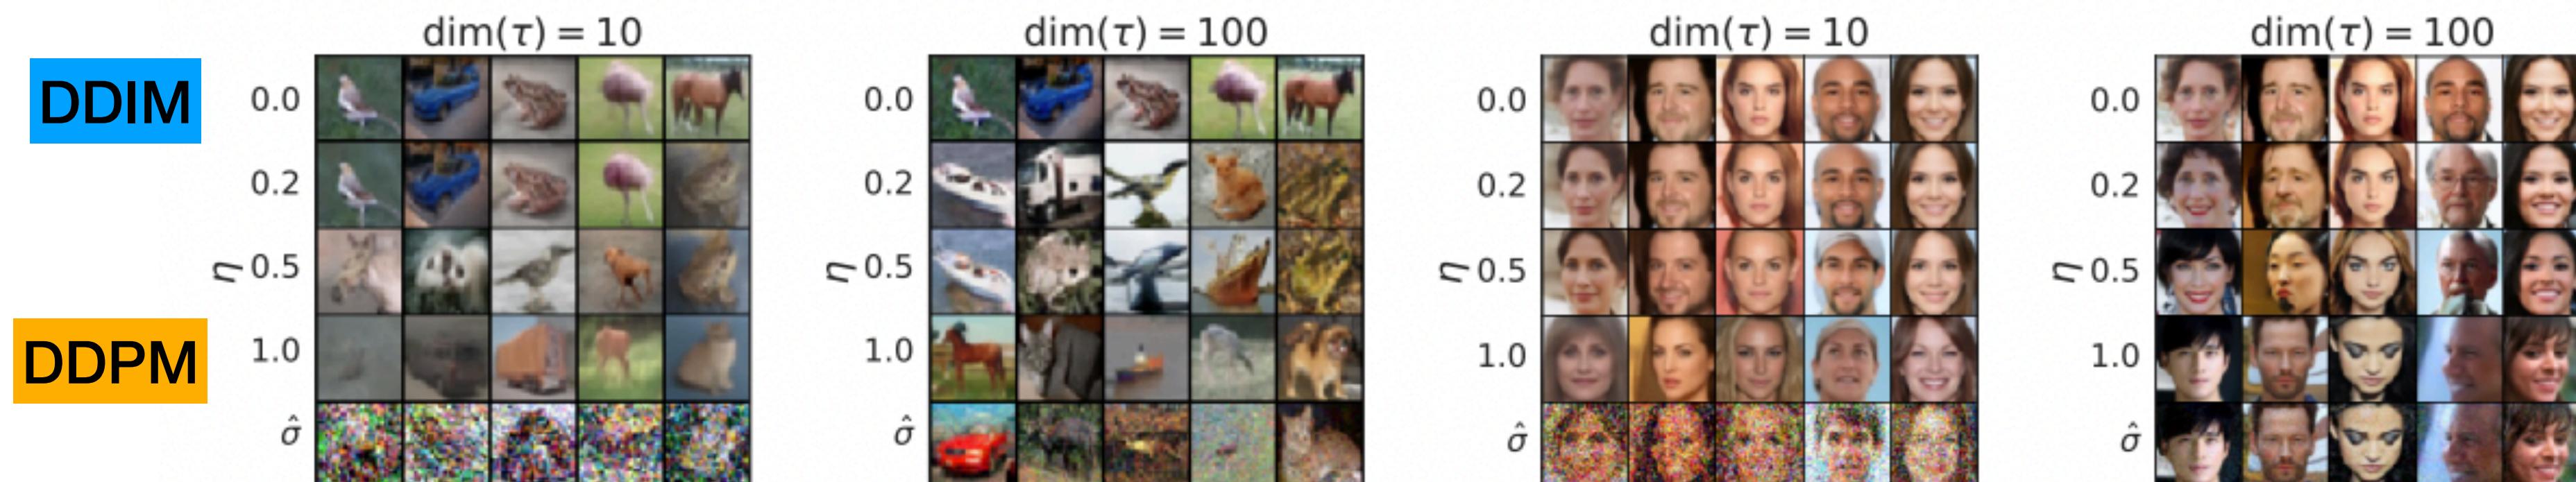
Generalized Posterior $q_\sigma(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}\left(\sqrt{\alpha_{t-1}}\mathbf{x}_0 + \sqrt{1 - \alpha_{t-1} - \sigma_t^2} \cdot \frac{\mathbf{x}_t - \sqrt{\alpha_t}\mathbf{x}_0}{\sqrt{1 - \alpha_t}}, \sigma_t^2 \mathbf{I}\right)$

$\sigma_t = 0$ 이면 DDIM

DDPM vs. DDIM

- $\sigma_{\tau_i}(\eta) = \eta \sqrt{(1 - \alpha_{\tau_{i-1}})/(1 - \alpha_{\tau_i})} \sqrt{1 - \alpha_{\tau_i}/\alpha_{\tau_{i-1}}}$ 로 두고 η 를 0.0, 0.2, 0.5, 1.0로 바꿔가면서 실험
($\eta = 0$ 이면 DDIM, $\eta = 1$ 이면 DDPM)
- 1000 steps로 트레이닝된 모델을 10, 20, 50, 100, 1000 steps로 inference해가면서 실험
- DDIM이 전체적으로 FID값이 더 낮으며 sampling step 수가 적을 때 더욱 차이가 난다

		CIFAR10 (32×32)					CelebA (64×64)				
		10	20	50	100	1000	10	20	50	100	1000
S		10	20	50	100	1000	10	20	50	100	1000
DDIM	0.0	13.36	6.84	4.67	4.16	4.04	17.33	13.73	9.17	6.53	3.51
	0.2	14.04	7.11	4.77	4.25	4.09	17.66	14.11	9.51	6.79	3.64
	0.5	16.66	8.35	5.25	4.46	4.29	19.86	16.06	11.01	8.09	4.28
DDPM	1.0	41.07	18.36	8.01	5.78	4.73	33.12	26.03	18.48	13.93	5.98
$\hat{\sigma}$		367.43	133.37	32.72	9.99	3.17	299.71	183.83	71.71	45.20	3.26



DDIM

- DDIM은 DDPM의 posterior를 일반화하고 stochastic한 항을 지우고 deterministic function으로 만든 것!
- 왜 Denoising Diffusion “Implicit” Models인가?

*“Implicit generative models use a latent variable z and transform it using **a deterministic function G_θ** that maps from $\mathbb{R}^m \rightarrow \mathbb{R}^n$ using parameters θ .”*

Mohamed, Shakir, and Balaji Lakshminarayanan. "Learning in implicit generative models." arXiv preprint arXiv:1610.03483 (2016).

DDIM

Keras Example

DDIM - def denoise

```
def denoise(self, noisy_images, noise_rates, signal_rates, training):
    # the exponential moving average weights are used at evaluation
    1. if training:
        network = self.network
    else:
        network = self.ema_network

    # predict noise component and calculate the image component using it
    2. pred_noises = network([noisy_images, noise_rates**2], training=training)
    pred_images = (noisy_images - noise_rates * pred_noises) / signal_rates

    return pred_noises, pred_images
```

<https://keras.io/examples/generative/ddim/>

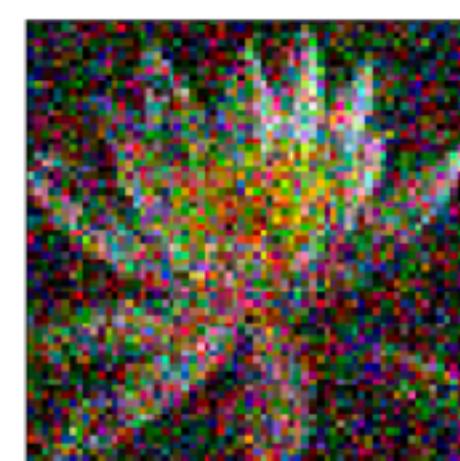
1. Training network와 EMA network를 별도로 두고, training시에는 training network를, inference시에는 EMA network를 사용

2. Network에 noisy image \mathbf{x}_t 를 넣고, predicted noise를 출력



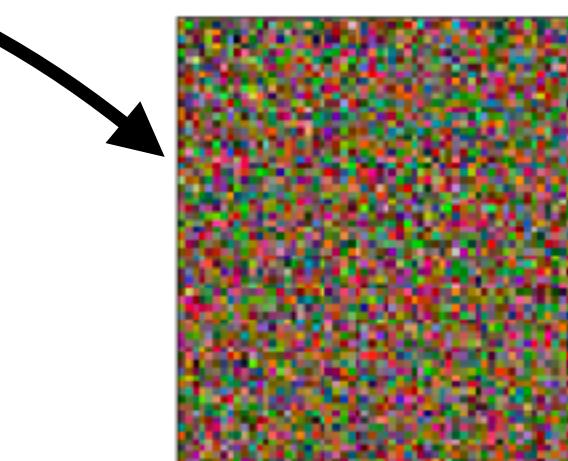
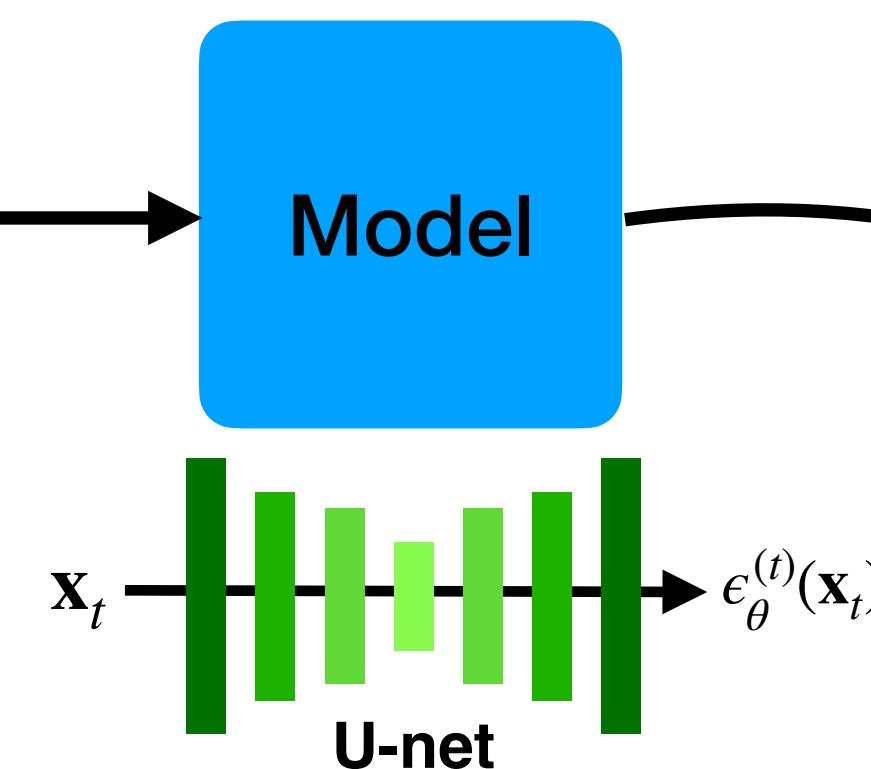
Predicted Image

$$\mathbf{x}_0 = \frac{\mathbf{x}_t - \sqrt{1 - \alpha_t} \epsilon_{\theta}^{(t)}(\mathbf{x}_t)}{\sqrt{\alpha_t}}$$



Noisy Image \mathbf{x}_t

$$\mathbf{x}_t = \sqrt{\alpha_t} \mathbf{x}_0 + \sqrt{1 - \alpha_t} \epsilon_t$$

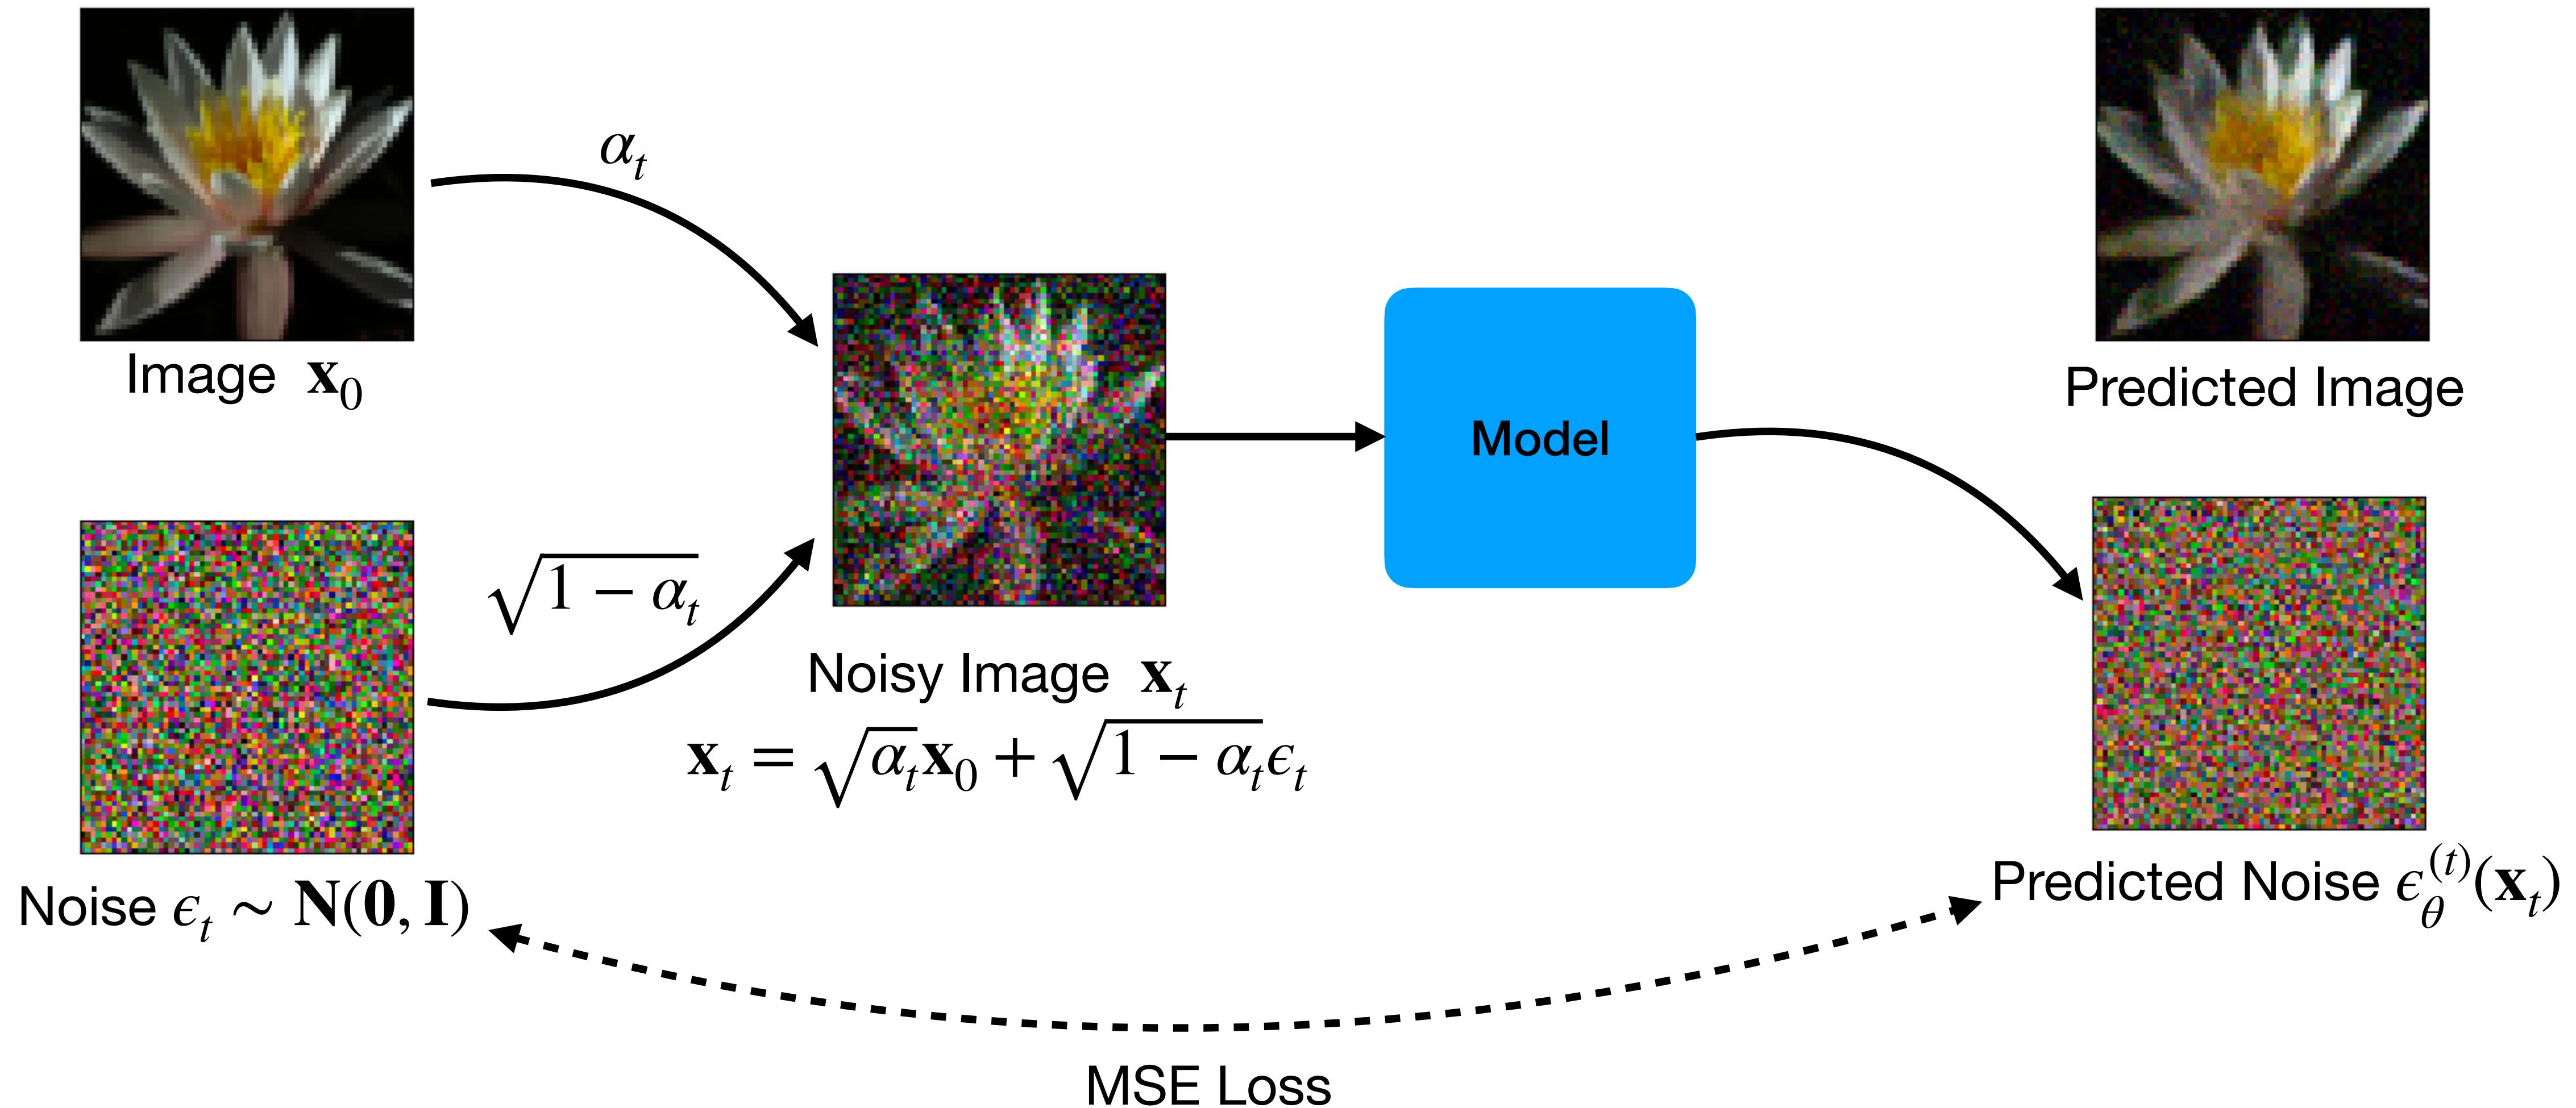


Predicted Noise

$$\epsilon_{\theta}^{(t)}(\mathbf{x}_t)$$

DDIM - def train_step

$$\mathbf{x}_0 = \frac{\mathbf{x}_t - \sqrt{1 - \alpha_t} \epsilon_{\theta}^{(t)}(\mathbf{x}_t)}{\sqrt{\alpha_t}}$$



DDIM - def train_step

```
def train_step(self, images):
    # normalize images to have standard deviation of 1, like the noises
    1. images = self.normalizer(images, training=True)
    noises = tf.random.normal(shape=(batch_size, image_size, image_size, 3))

    2. # sample uniform random diffusion times
    diffusion_times = tf.random.uniform(
        shape=(batch_size, 1, 1, 1), minval=0.0, maxval=1.0
    )
    noise_rates, signal_rates = self.diffusion_schedule(diffusion_times)

    3. # mix the images with noises accordingly
    noisy_images = signal_rates * images + noise_rates * noises

    with tf.GradientTape() as tape:
        # train the network to separate noisy images to their components
        5. pred_noises, pred_images = self.denoise(
            noisy_images, noise_rates, signal_rates, training=True
        )

        noise_loss = self.loss(noises, pred_noises) # used for training
        image_loss = self.loss(images, pred_images) # only used as metric

    6. gradients = tape.gradient(noise_loss, self.network.trainable_weights)
    self.optimizer.apply_gradients(zip(gradients, self.network.trainable_weights))

    self.noise_loss_tracker.update_state(noise_loss)
    self.image_loss_tracker.update_state(image_loss)

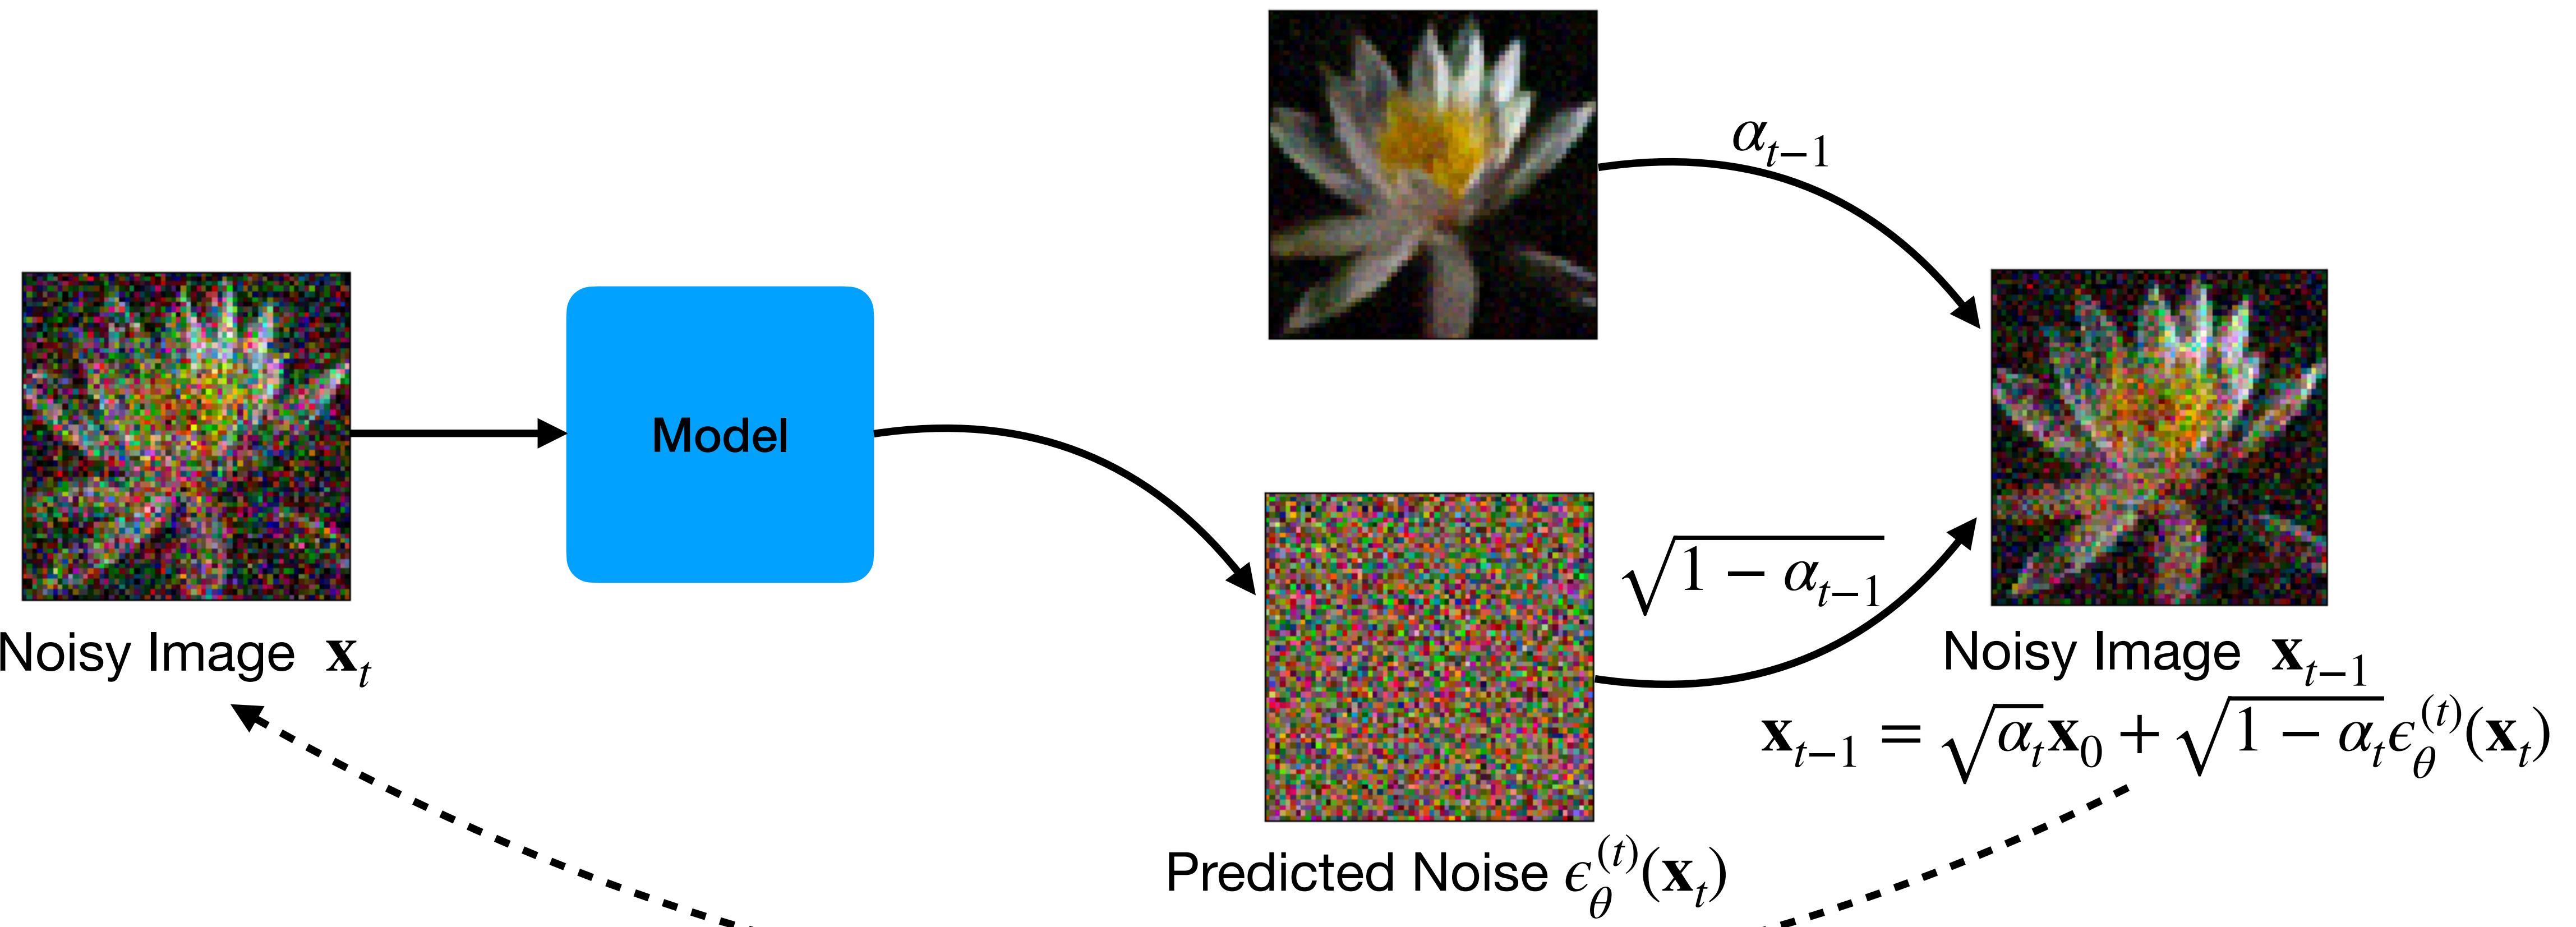
    8. # track the exponential moving averages of weights
    for weight, ema_weight in zip(self.network.weights, self.ema_network.weights):
        ema_weight.assign(ema * ema_weight + (1 - ema) * weight)

    # KID is not measured during the training phase for computational efficiency
    return {m.name: m.result() for m in self.metrics[:-1]}
```

1. Image \mathbf{x}_0 를 normalization하고, noise $\epsilon_t \sim \mathbf{N}(\mathbf{0}, \mathbf{I})$ 를 샘플링
2. Diffusion time $t \sim U(0,1)$ 를 샘플링
3. Diffusion schedule에 따라 signal rate $\sqrt{\alpha_t}$ 와 noise rate $\sqrt{1 - \alpha_t}$ 를 구함
4. Noisy image \mathbf{x}_t 를 만듬. $\mathbf{x}_t = \sqrt{\alpha_t} \mathbf{x}_0 + \sqrt{1 - \alpha_t} \epsilon_t$
5. Model에 noisy image \mathbf{x}_t 를 넣어 predicted image와 predicted noise $\epsilon_{\theta}^{(t)}(\mathbf{x}_t)$ 출력
6. 본래의 noise ϵ 와 predicted noise $\epsilon_{\theta}^{(t)}(\mathbf{x}_t)$ 간에 MSE loss 잡음
7. Gradient를 구하고 optimizer를 이용하여 weight 업데이트
8. EMA (Exponential Moving Average) model의 weight를 업데이트

DDIM - def reverse_diffusion

$$\mathbf{x}_0 = \frac{\mathbf{x}_t - \sqrt{1 - \alpha_t} \epsilon_{\theta}^{(t)}(\mathbf{x}_t)}{\sqrt{\alpha_t}}$$



DDIM - def reverse_diffusion

```
def reverse_diffusion(self, initial_noise, diffusion_steps):
    # reverse diffusion = sampling
    num_images = initial_noise.shape[0]
    1. step_size = 1.0 / diffusion_steps

    # important line:
    # at the first sampling step, the "noisy image" is pure noise
    # but its signal rate is assumed to be nonzero (min_signal_rate)
    next_noisy_images = initial_noise
    2. for step in range(diffusion_steps):
        noisy_images = next_noisy_images

        # separate the current noisy image to its components
        diffusion_times = tf.ones((num_images, 1, 1, 1)) - step * step_size
        noise_rates, signal_rates = self.diffusion_schedule(diffusion_times)
        pred_noises, pred_images = self.denoise(
            noisy_images, noise_rates, signal_rates, training=False
        )
        # network used in eval mode

        # remix the predicted components using the next signal and noise rates
        5. next_diffusion_times = diffusion_times - step_size
        next_noise_rates, next_signal_rates = self.diffusion_schedule(
            next_diffusion_times
        )
        next_noisy_images = (
            next_signal_rates * pred_images + next_noise_rates * pred_noises
        )
        # this new noisy image will be used in the next step

    return pred_images
```

1. Batch size인 num_images 구하고, sampling을 수행할 step 수 diffusion_steps으로부터 step_size 구함
2. initial_noise에서부터 생성 iteration 시작
3. diffusion_times는 1에서부터 0까지 step_size만큼씩 줄여감, Diffusion schedule에 따라 signal rate $\sqrt{\alpha_t}$ 와 noise rate $\sqrt{1 - \alpha_t}$ 를 구함
4. Model에 noisy_images \mathbf{x}_t 를 넣어 predicted image \mathbf{x}_0 와 predicted noise $\epsilon_{\theta}^{(t)}(\mathbf{x}_t)$ 출력
5. next_diffusion_times을 구하고, diffusion schedule에 따라 signal rate $\sqrt{\alpha_{t-1}}$ 와 noise rate $\sqrt{1 - \alpha_t}$ 를 구함
6. 다음 step의 Noisy image \mathbf{x}_{t-1} 를 만듬.
$$\mathbf{x}_{t-1} = \sqrt{\alpha_{t-1}} \mathbf{x}_0 + \sqrt{1 - \alpha_{t-1}} \epsilon_{\theta}^{(t)}(\mathbf{x}_t)$$